



TECHNISCHE HOCHSCHULE NÜRNBERG  
GEORG SIMON OHM

Fakultät Informatik

# Enablement of Kubernetes Based Open-Source Projects on IBM Z

Bachelorarbeit im Studiengang Informatik

vorgelegt von

Sarah Julia Kriesch

Matrikelnummer 303 6764

Erstgutachter:	Prof. Dr. Ralf-Ulrich Kern
Zweitgutachter:	Prof. Dr. Tobias Bocklet
Betreuer:	M.Sc. Alice Frosi
Unternehmen:	IBM Deutschland R & D GmbH

© 2020

Dieses Werk einschließlich seiner Teile ist **urheberrechtlich geschützt**. Jede Verwertung außerhalb der engen Grenzen des Urheberrechtsgesetzes ist ohne Zustimmung des Autors unzulässig und strafbar. Das gilt insbesondere für Vervielfältigungen, Übersetzungen, Mikroverfilmungen sowie die Einspeicherung und Verarbeitung in elektronischen Systemen.

Angaben des bzw. der Studierenden:

Name: \_\_\_\_\_ Vorname: \_\_\_\_\_ Matrikel-Nr.: \_\_\_\_\_

Fakultät: Studiengang:

Semester:

**Titel der Abschlussarbeit:**

Ich versichere, dass ich die Arbeit selbständig verfasst, nicht anderweitig für Prüfungszwecke vorgelegt, alle benutzten Quellen und Hilfsmittel angegeben sowie wörtliche und sinngemäße Zitate als solche gekennzeichnet habe.

Ort, Datum, Unterschrift Studierende/Studierender

### Erklärung zur Veröffentlichung der vorstehend bezeichneten Abschlussarbeit

Die Entscheidung über die vollständige oder auszugsweise Veröffentlichung der Abschlussarbeit liegt grundsätzlich erst einmal allein in der Zuständigkeit der/des studentischen Verfasserin/Verfassers. Nach dem Urheberrechtsgesetz (UrhG) erwirbt die Verfasserin/der Verfasser einer Abschlussarbeit mit Anfertigung ihrer/seiner Arbeit das alleinige Urheberrecht und grundsätzlich auch die hieraus resultierenden Nutzungsrechte wie z.B. Erstveröffentlichung (§ 12 UrhG), Verbreitung (§ 17 UrhG), Vervielfältigung (§ 16 UrhG), Online-Nutzung usw., also alle Rechte, die die nicht-kommerzielle oder kommerzielle Verwertung betreffen.

Die Hochschule und deren Beschäftigte werden Abschlussarbeiten oder Teile davon nicht ohne Zustimmung der/des studentischen Verfasserin/Verfassers veröffentlichen, insbesondere nicht öffentlich zugänglich in die Bibliothek der Hochschule einstellen.

Hiermit ☐ genehmige ich, wenn und soweit keine entgegenstehenden Vereinbarungen mit Dritten getroffen worden sind,

☐ genehmige ich nicht,

dass die oben genannte Abschlussarbeit durch die Technische Hochschule Nürnberg Georg Simon Ohm, ggf. nach Ablauf einer mittels eines auf der Abschlussarbeit aufgebrachten Sperrvermerks kenntlich gemachten Sperrfrist

von                      Jahren (0 - 5 Jahren ab Datum der Abgabe der Arbeit),

der Öffentlichkeit zugänglich gemacht wird. Im Falle der Genehmigung erfolgt diese unwiderruflich; hierzu wird der Abschlussarbeit ein Exemplar im digitalisierten PDF-Format auf einem Datenträger beigelegt. Bestimmungen der jeweils geltenden Studien- und Prüfungsordnung über Art und Umfang der im Rahmen der Arbeit abzugebenden Exemplare und Materialien werden hierdurch nicht berührt.

---

Ort, Datum, Unterschrift Studierende/Studierender

## Kurzdarstellung

Open-Source-Projekte entwickeln Software, die frei verfügbar ist. Diese Communitys lassen automatisierte Tests gegen jede Code-Änderung laufen, um die bester Software-Qualität zu garantieren. Diese Tests sollen auf unterschiedlichen Architekturen laufen können. Es ist schwierig, Software für grundlegende Hardware ohne Zugriff darauf zu testen. Deshalb muss die Architektur s390x für Z Systeme für ausgewählte Open-Source-Projekte emuliert (nachgeahmt) und in die entsprechende CI/CD-Pipeline eingefügt werden. Das sollte mit schnellen Deployment-Methoden im Emulator QEMU durchgeführt werden. Kubernetes wird als containerisiertes Beispiel-Projekt aus der Sicht als Grundlage für dafür eingeführte Anwendungen verwendet. Ein weiteres Open-Source-Projekt, Apache Cassandra, wird verwendet um Tests auf der Anwendungsschicht im Kubernetes-Stack zu repräsentieren. Zusätzlich müssen die minimalen Systemanforderungen für die Einrichtung innerhalb der CI/CD-Infrastruktur beider Projekte wegen der Minimierung an Festplattenplatz, Memory und CPU-Verbrauch analysiert werden. Zum Abschluss wird die automatische Emulation beider Projekte in die Test-Infrastruktur integriert, so dass diese Projekte für die Architektur von Z Systemen aktiviert sind. Allgemein soll diese Methode für weitere Open-Source-Projekte in der Zukunft wiederverwendet werden können.

## Abstract

Open-source-projects are developing software which is freely available. These communities are running automated tests against every code change in order to guarantee the best software quality. These tests should be able to run on different architectures. It is difficult to test software for essential hardware without access. Therefore, the architecture s390x for Z systems has to be emulated on x86 for chosen open-source projects and included in their CI/CD pipeline. That should be executed with fast deployment methods in the emulator QEMU. Kubernetes is used as a containerized example project in the point of view as the foundation for instituting applications. Another open-source project, Apache Cassandra, is applied to represent tests on the application layer in the Kubernetes stack. Additionally, minimal system requirements have to be analyzed for the setup inside of the CI/CD infrastructure of both projects concerning the minimization of space, memory and CPU usage for deployments. Finally, the automated emulation of both projects will be integrated into the test infrastructure, that these projects are enabled for the architecture of Z systems. Overall, this method should be reapplied for further open-source-projects in the future.



# Contents

<b>1. Introduction</b>	<b>1</b>
1.1. Mainframe Computers	2
1.2. Hardware Emulation	2
1.3. Open Source Projects	3
1.3.1. Docker	3
1.3.2. Kubernetes	4
1.3.3. Apache Cassandra	4
<b>2. Emulation</b>	<b>7</b>
2.1. QEMU	7
2.2. Full-System Emulation	8
2.3. User-Mode Emulation	8
2.4. Emulation of Alternative Architectures	9
2.4.1. Binfmt_misc	10
2.4.2. Qemu-user-static	10
2.4.3. BuildX	11
2.4.4. Multi-Architecture Images	11
2.5. File Systems	13
2.5.1. UnionFS	13
2.5.2. Ext4	14
2.5.3. Mounting an External File System	15
<b>3. Prerequisites</b>	<b>17</b>
3.1. Registration of Qemu-S390x	17
3.2. Enabling Multi-Architecture Images	18
3.3. Fetching Signed Linux Kernel Image	19
3.4. Optimized QEMU Command	19
<b>4. Continuous Integration</b>	<b>21</b>
4.1. CI/CD in the Software Development Lifecycle	21
4.2. CI/CD Systems	21
<b>5. Kubernetes</b>	<b>23</b>
5.1. Overview	23

5.2. Multi-Staging Dockerfile . . . . .	23
5.3. Installation . . . . .	24
5.4. Installation of the Ultimate Go . . . . .	25
5.5. Building Tests . . . . .	26
5.5.1. E2e.test . . . . .	27
5.6. Run in the Main Dockerfile . . . . .	27
5.7. Integration into CI/CD . . . . .	27
<b>6. Cassandra . . . . .</b>	<b>29</b>
6.1. Overview . . . . .	29
6.2. Installation . . . . .	29
6.3. Java Optimization . . . . .	30
6.4. Nagios Monitoring Check for Tests . . . . .	31
6.5. Workaround Because of an JVM Issue . . . . .	32
6.6. Run Docker Image in QEMU . . . . .	34
6.7. Integration into CI/CD . . . . .	35
<b>7. Outlook . . . . .</b>	<b>37</b>
<b>8. Summary . . . . .</b>	<b>39</b>
<b>A. Supplemental Information . . . . .</b>	<b>41</b>
A.1. Circle CI configuration for Apache Cassandra . . . . .	42
<b>List of Figures . . . . .</b>	<b>45</b>
<b>List of Tables . . . . .</b>	<b>47</b>
<b>Bibliography . . . . .</b>	<b>49</b>
<b>Glossary . . . . .</b>	<b>53</b>

# Chapter 1.

## Introduction

One business introduced by IBM is the mainframe, now known as a Z system. It is possible to run Linux on it. There is a large community behind Linux and open-source projects. Open source does not contain only Linux. There are different applications and other software developed by open-source communities. The mainframe hardware architecture is different from the architecture of a home PC. The Z system architecture is called s390x and a default system x86. Not every open-source community has got access to such a Z system. Most open-source contributors have got systems with the architecture x86. Therefore, it should be possible to test hardware dependencies for s390x on available systems. The goal of this Bachelor Thesis is to integrate emulated Z systems for different open-source projects to test their software for hardware dependencies so that it allows to release new versions in the CI/CD pipeline of the respective project for running on the architecture s390x without the available hardware. Deployments of the latest software version on Github and running tests have to be automated for that.

Our focus is on Kubernetes-based open-source projects. These should be emulated for Z systems in the CI/CD test infrastructure by these chosen projects. That will be done on systems by these communities. As the first step, the emulator will be chosen with the focus on functionality for Z systems on x86 architecture. Afterwards, Kubernetes is installed in a Docker container. Additionally, tests should be able to be run on this system. That all will be integrated into the emulation environment for an automated start. In conclusion, the CI/CD system should be able to execute all tests.

The same will be done with the NoSQL database Cassandra for the Apache community to represent the whole system stack from Kubernetes until the application layer for container platforms. Other points are minimal systems requirements and minimal systems sizes. Here are different methods evaluated to minimize the system for emulation.

## 1.1. Mainframe Computers

Mainframe computers are large computers. Some of these computers are part of the Z series<sup>1</sup> by IBM. They are not only used as internet servers, but also for mission-critical apps and blockchain. Mainframes can handle large numbers of transactions in one second [Tane 14, p.56]. Thousands of virtual machines can run on such a system. Z systems do not use the well known x86 architecture. They are built with s390x. The current architecture version has been introduced in late 2000 by IBM [Bloc 19, p.15]. The mainframe has built in the fastest available processor with 5 GHz with on core and on chip cache, extensive memory and dedicated I/O processing [Linu].

IBM Z contains "IBM Z pervasive encryption" for comprehensive protection around the data on the system [Lasc 20, p.4]. Such systems are offering a horizontal and upright scaling of processes, which allows the operation of many hundred virtual systems in parallel [Tsch 09, p.13]. The traditional operating system for mainframes has been z/OS. Z systems have been optimized for open-source software as Linux [Lasc 20, p.8]. The goal by IBM has been to offer a combination of a robust and securable hardware platform with the power of different Linux distributions. Ubuntu Linux is used as a base operating system for this Bachelor Thesis.

## 1.2. Hardware Emulation

Not everybody has access to hardware with essential architecture. The software should be able to run on most relevant hardware architectures. The solution for Software Developers is hardware emulation. You can test with the hardware emulation whether the software is running correctly. Accordingly, you can run different operating systems and applications for specialized hardware in emulators or virtualization software. It is possible to enable other hardware architectures than the one of the host. That will be done with the implementation of a guest system on a host with a different target architecture [Rose 15, p.3]. In our case, it should be possible to run applications for mainframes with the target architecture s390x on a host system with the architecture x86.

---

<sup>1</sup><https://www.ibm.com/it-infrastructure/z/hardware/>



## 1.3. Open Source Projects

### 1.3.1. Docker

Docker<sup>2</sup> is one possible container engine used by Kubernetes. This technology emerged in 2013 as a base for future container technologies. The difference to existing container technologies (LXC as an example) has been the possibility of distributing systems. Docker Inc. has established a public container registry with the name Docker Hub for public usable Docker images. That should facilitate setups and the entrance into work with containers. Another benefit of Docker is that all installation and configuration steps for a system are described in one reusable text file. This file (Dockerfile) is the foundation of Docker images and most public Docker images on Docker Hub are maintained. Docker images are allocatable systems obtained from a container registry with `docker pull`, or built based on a self-written Dockerfile for the local registry on the server with `docker build -t ${name} ..`

`-t` is the tag and defines the name of the docker image listed in the registry with the command `docker images`. Only successful builds of images can be registered in such a container registry.

The **Dockerfile** contains different instructions for building steps. The line with the "FROM" command defines the base image. Every public or local usable Docker image can be used as a base image with the whole operating system incl. pre-installed packages and relevant configuration. Docker images include only a minimal operating system defined in the Dockerfile of the base image together with all listed packages for installation. These installations are listed in the Dockerfile with the command "RUN" before apt, yum, pip or other installation commands. After the build, all commands are registered in a Docker manifest within JSON format together containing all information about the Docker image. If a new application should be executed inside of the Docker container, then creating a directory for this application besides the Dockerfile is possible. The command "ADD" can integrate this application into the Docker container during the build process. Such a line has the following structure:

```
ADD ./dir/app.py /app.py
```

This application can be started with the "CMD" command then:

```
CMD ["python", "/app.py"]
```

It is recommended to define one service or process for one single container and to connect all containers for a start then. One single container is launchable with the command `docker run ${name}`. In this case, `${name}` can be the name of the image or the image id. The command `docker-compose` is available to automate the setup with multiple connected containers.

The Dockerfile format has been spread as an easily understandable base for most container runtimes. Docker has turned out to be more like a developer tool for many container

---

<sup>2</sup><https://www.docker.com/>

technologies in the last years. Consequently, as an example, Kubernetes is using Docker for their community tests as a base container runtime. Most new container orchestration platforms are developed based on Docker. The difference is the replacement of the docker command and the expansion with additional features for clustering and specialized configurations.

### 1.3.2. Kubernetes

Kubernetes<sup>3</sup> is an open-source project for container orchestration, well known as k8s. Google started this project. A Kubernetes cluster has at least one Master node and one Worker node for high availability. One node is one host. The Master node is responsible for managing all Worker nodes with applications. All configurations and distributions are performed from there. New Worker nodes are added to the Kubernetes cluster with "join" on the Master node, too. The Worker node is deploying containers with different services. Every Worker node can communicate with other Worker nodes in the cluster.

Kubernetes is scalable for using containers as distributed services in a pod. A pod is representing something as a single server split into different connected containers with all services for a running application. Pods can be replicated to multiple nodes for high availability. Kubernetes is configurable with different container runtimes, like Docker, containerd<sup>4</sup> or CRI-O<sup>5</sup>, for example. The Container Runtime Interface (CRI) is necessary for managing container images, the life cycle of container pods, networking and help functions [Scho 19, p.16]. The most used container runtime is Docker in the Kubernetes project. The filesystem of a container is described in the format of a Dockerfile. Most container runtimes support this format.

### 1.3.3. Apache Cassandra

Apache Cassandra<sup>6</sup> is a NoSQL database management system developed by the Apache Foundation. The project had been started internally at Facebook and has been released as an open-source project in 2008. Cassandra provides continuous availability, high performance, and linear scalability besides offering operational simplicity and effortless replication across data centres and geographies [Data]. It is recommended for mission-critical data. The Cassandra Query Language (CQL) is similar to SQL and includes JSON support. The similarity simplifies migrations from relational database management systems to Apache Cassandra.

CQL includes an abstraction layer that hides implementation details of the structure.

NoSQL database stores are using other methods than relational database management

---

<sup>3</sup><https://kubernetes.io/>

<sup>4</sup><https://containerd.io/>

<sup>5</sup><https://cri-o.io/>

<sup>6</sup><https://cassandra.apache.org/>

systems for data access and own a different structure. Cassandra is using hashing to fetch rows from the table as a column-oriented database management system (DBMS). Such a DBMS stores data tables by column rather than by row.

NoSQL databases were developed as massively scalable database management systems that can write and read data anywhere while distributing availability to billions of users. This field is a part of Big Data.



## Chapter 2.

### Emulation

#### 2.1. QEMU

QEMU is an open-source emulator introduced by Fabrice Bellard in the year 2003 and available in most Linux distributions now. It is embedded in different virtualization tools as KVM and XEN, too. QEMU is well tested and contains all the necessary features for emulations of other architectures on alternative hardware. QEMU is a generic emulator for different system architectures. It can also be used for emulation of obsolete hardware [Opsa 13, p.24]. It has been extended for various architectures as x86, ARM, PowerPC, Sparc32, Sparc64, MIPS and s390x. It is a processor emulator using binary translation [Butt 11] which is executing and translating emulated instructions based on blocks. Each block comprises one entry and one exit point [Wang 10, p.5]. The binary translation will be executed with the "Tiny Code Generator" (TCG) inside of QEMU. The TCG is a small compiler replacing GCC because of unlimited releases and code changes. The TCG is converting the blocks of target instructions into a standardized form. Subsequently, it has to be compiled for the host or target architecture. If a binary for a new target architecture is necessary, the frontend of TCG will be converted while QEMU is ported to a new architecture. The TCG integrates new code for the new host architecture in the background then. It is also dedicated to improving performance with avoiding repeated translations by buffering already translated code [Cota 17]. The TCG takes care of the emulation of the guest processor running as a thread launched by QEMU. The memory of a guest is allocated during launch. Then that is mapped into the address space of the QEMU process [Opsa 13, p.29]. It is possible to run unmodified guest operating systems. The open-source projects can use any Linux distribution as their base operating system then because QEMU is integrated as a default package. QEMU does not emulate the entire hardware. That is only possible for the CPU. QEMU is used for emulations in this Bachelor Thesis.

## 2.2. Full-System Emulation

The full-system emulation emulates a whole system with hardware, the operating system (with the Linux kernel) and the userspace (with application processes). The system (hardware and operating system) will be translated unmodified. The condition of system emulation (compared to user mode emulation) is that you can run privileged instructions [Butt 11, p.2]. This feature enables the translation of the unmodified target code on the operating system. That all leads to a slowdown of the emulation in comparison to the user-mode emulation. Additionally, this emulation can be used as an application development platform where specific hardware is not available. System emulation with emulated hardware is slower than a real machine because instructions should be executed directly in the guest hardware. However, that has to be emulated in software. That implies multiple host instructions as a result because of the translation for a single guest instruction [Tong 14, p.1]. It is possible to reduce the supported and attached additional devices with additional options. The deactivation of a graphic card can be specified with **-nographic** as an example. Afterwards, less hardware has to be emulated, which has given a better performance.

System emulation benefits from the virtualization support as KVM. In this case, CPU operations are not required to emulate.

## 2.3. User-Mode Emulation

The user-mode emulation does not emulate the whole system. It is faster than the full-system emulation because it does not engage so much hardware resources. Application processes can be run in QEMU with a minimal system for a specific application. This emulation type is working on a system call level. Therefore, the application has to be runnable as a single process executable itself. The emulator is using the Linux kernel to emulate system signal calls then. That can be managed by mapping system calls of the target system to an equivalent system call on the host with threading (with a separate virtual CPU) [QEMU]. This process does not require the emulation of the full memory management unit [Butt 11, p.2]. The user-mode emulation can run directly non-privileged instructions or is using system calls to ask for a selected service from the operating system.

The disadvantage of user-mode emulation is the only possibility to run single processes. Most services contain different applications with the result of multiple processes. Consequently, these applications need a full system emulation.

## 2.4. Emulation of Alternative Architectures

It is achievable to emulate alternative architectures on another hardware architecture. The package `qemu-user-static` has to be installed then, and the chosen architecture has to be registered in `binfmt_misc`. `binfmt_misc` is a kernel module. You can register other architectures within that, so that multiple other architectures can be run on one host. Not only QEMU can be used for system emulation. But also, container technologies as Docker include an integrated QEMU compatibility as an additional emulation feature for building images for other architectures. This technology has got the name [2.4.3"BuildX"](#). Accordingly, a hybrid virtualization approach is practicable with different virtualization and emulation technologies as with QEMU and Docker together. In this case, an external Linux kernel will be integrated into QEMU, and the application can be mounted via a loaded Docker image in a hard disk image. The Linux kernel is required for QEMU and should be built or downloaded (see [3.3"Fetching Linux Kernel Image"](#)).

The `initrd` should match the Linux kernel and is optional. It is used to start an `init` process together with the required test scripts.

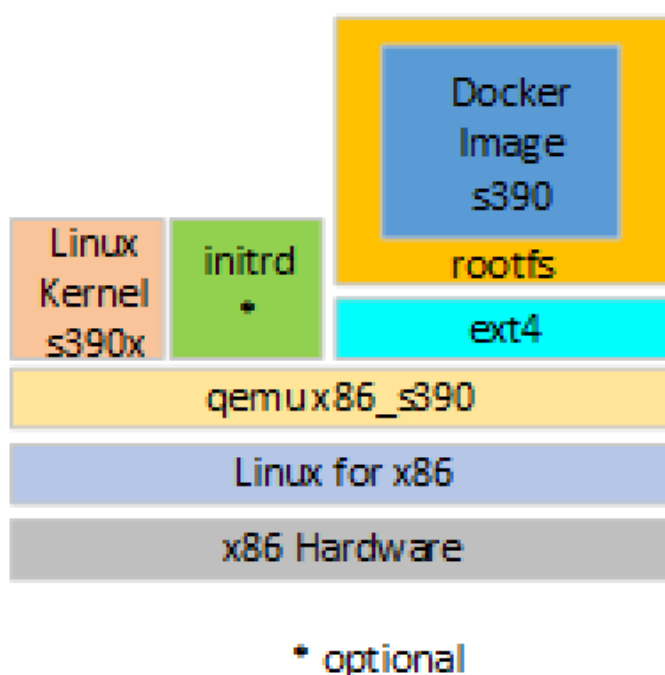


Figure 2.1.: Emulation with QEMU

### 2.4.1. Binfmt\_misc

Binfmt\_misc is a kernel module which allows invoking almost every program by simply typing its name in the shell. That permits to execute user space applications such as emulators and virtual machines for other hardware architectures. It recognizes the binary-type by matching some bytes at the beginning of the file with a magic number. These executable formats have to be registered in the file `/proc/sys/fs/binfmt_misc/register`.

The structure of the configuration for the registration is the following:

**:name:type:offset:magic:mask:interpreter:flags**

The **name** is the name of the architecture for the binary format. The **type** can be **E**, or **M**. E is for executable file formats as `.exe` for example. M is used for a format identified by a **magic** number at an absolute **offset** in the file, and the **mask** is a bitmask indicating which bits in the number are significant[Slac] (which is utilized in our case). The offset can be kept empty. The magic is a byte sequence of hexadecimal numbers. The **interpreter** is the program that should be invoked with the binary[Gün]. The path has to be specified for it. The field **flag** is optional. It checks different aspects of the invocation of the interpreter. The **F** flag is necessary in our case for "fix binary". That supposes that a new binary has to be spawned if the misc file format has been invoked.

In this way, it is possible to register another system architecture (s390x on x86) on the host system, and it is realizable to emulate the other architecture afterwards.

The necessary command for s390x registration can be found under 3.1. The package **binfmt-support** has to be installed first for using this kernel feature.

### 2.4.2. Qemu-user-static

The package **qemu-user-static**<sup>1</sup> enables the execution of multi-architecture container emulation based on **binfmt\_misc** with **QEMU**. This package contained in various Linux distributions includes a set of binfmt configurations for various architectures together with an amount of statically compiled QEMU emulators that alternative architectures can run on another architecture [Yang 19]. When building for one new architecture, static binaries are relevant because of the possibility of an "exec init error" without them.

The installation contains emulators for all available architectures supported by the QEMU project incl. s390x aligned with the specific host architecture. In this manner, you can build and run a binary for different architectures on the same host (see 3.4 Optimized Qemu Command). That is the base for 2.4.3 BuildX by Docker, too.

---

<sup>1</sup><https://github.com/multiarch/qemu-user-static>



### 2.4.3. BuildX

As mentioned above, Docker is using QEMU for emulations with BuildX<sup>2</sup>. That is an integrated "experimental Feature" since version Docker 19.03. That implies that it is not enabled as default because it is very new and for experiments. It has to be activated with `DOCKER_CLI_EXPERIMENTAL=enabled` (see 3.1 Registration of Qemu-S390x). The disadvantage of the provided BuildX inside of Docker is that you do not receive the up-to-date version of BuildX. There is a version with a more detailed output and better working with an additional **buildx** behind **docker build**. This version can be cloned from <https://github.com/docker/buildx> and installed with **make install**.

In general, BuildX is a Docker CLI plugin for expanding the "docker build" command. That includes the possibility of multi-architecture builds and exceptional output configuration. That involves not only the upload possibility to your local docker registry (see the output of **docker images**). That includes the output of "docker build" (which creates a docker image based on the requirements and installations in the used Dockerfile) into a local directory, a tar archive or a public registry on Docker Hub. Docker Hub is the most used and maintained registry for Docker images.

The **Dockerfile** is using the requirements by the base docker image in the FROM command on top of the Dockerfile together with installations and configurations executed by all RUN commands or attached scripts with the command ADD. The **-t** flag is adding a special name as a tag under **docker images**. The special architecture for multi-architecture<sup>2.4.4</sup> builds can be specified with the additional option **--platform**. The option **build create** provides the option to create different build instances for multiple build combinations of architectures.

### 2.4.4. Multi-Architecture Images

Default Linux distribution providers have to package their software for different system architectures because of different required drivers included in the Linux kernel and dependencies to these. That is the reason that Docker images have to be built equally for different architectures. That can be performed with different base images for the respective architectures or with multi-arch images. It is possible to use one Dockerfile for different architectures then. One example for such a multi-arch Docker image is adoptopenjdk/adoptopenjdk8 based on Ubuntu<sup>3</sup>. In this case hardware dependencies are selected with `dpkg --print-architecture` and included as ARCH into the different cases for downloading the relevant tar archive of OpenJDK:

---

<sup>2</sup><https://github.com/docker/buildx>

<sup>3</sup><https://github.com/AdoptOpenJDK/openjdk-docker/blob/master/8/jdk/ubuntu/Dockerfile.hotspot.nightly.full>

```

RUN set -eux; \
ARCH="$(dpkg --print-architecture)"; \
case "${ARCH}" in \
    aarch64|arm64) \
        ESUM='2c6540ff8ea3d89362fd02143b24303e2359be249a6be6cb7e6580472686d863'; \
        BINARY_URL='https://github.com/AdoptOpenJDK/openjdk8-binaries/releases/download/jdk8u-2020-08-05-08-20/OpenJDK8U-jdk_aarch64_linux_hotspot_2020-08-05-08-20.tar.gz'; \
        ;; \
    armhf|armv7l) \
        ESUM='3c57c121415ef721ba9c73dfa809cd2f689d7369ef3d92d055f7c7c81ed2d697'; \
        BINARY_URL='https://github.com/AdoptOpenJDK/openjdk8-binaries/releases/download/jdk8u-2020-07-31-04-43/OpenJDK8U-jdk_arm_linux_hotspot_2020-07-31-04-43.tar.gz'; \
        ;; \
    ppc64el|ppc64le) \
        ESUM='668972e8d6d0815c03b32dac3f5c663453b4b0842938d13cfb0c1232a26a8087'; \
        BINARY_URL='https://github.com/AdoptOpenJDK/openjdk8-binaries/releases/download/jdk8u-2020-08-05-08-20/OpenJDK8U-jdk_ppc64le_linux_hotspot_2020-08-05-08-20.tar.gz'; \
        ;; \
    s390x) \
        ESUM='045aa287c48fd84eef19f2828ed1d15d27059a9328f5ca2cc91f6d42489af956'; \
        BINARY_URL='https://github.com/AdoptOpenJDK/openjdk8-binaries/releases/download/jdk8u-2020-08-05-08-20/OpenJDK8U-jdk_s390x_linux_hotspot_2020-08-05-08-20.tar.gz'; \
        ;; \
    amd64|x86_64) \
        ESUM='c26124b14c6a42d89278d3ce108b43c1210a317aba0163d985c06a79a695a174'; \
        BINARY_URL='https://github.com/AdoptOpenJDK/openjdk8-binaries/releases/download/jdk8u-2020-08-05-08-20/OpenJDK8U-jdk_x64_linux_hotspot_2020-08-05-08-20.tar.gz'; \
        ;; \
    *) \
        echo "Unsupported arch: ${ARCH}"; \
        exit 1; \
        ;; \
esac;

```

Figure 2.2.: Multi-Arch Image AdoptOpenJDK 8

You can see all URLs for arm64, armv7l, ppc64le, s390x and x86 in this case. This RUN command in the Dockerfile is using the architecture given by the option

```
--build-arg ARCH=s390x
```

as an example in the `docker build` command or the `--platform` option will be used with BuildX. `Docker build` is stacked against BuildX, because only one image can be built with one command. Buildx provides the possibility to consign different architectures (comma separated) to the platform option. Thereby, one command can build multiple images for different architectures.

## 2.5. File Systems

Docker is using the Union File System which is not compatible with QEMU.

QEMU can integrate only hard disk formats for default Linux file systems as ext2, ext3, ext4, XFS and Btrfs. The driver `virtio_blk` is used to mount external file systems and emulates read/write in the physical block device[Barb 18]. Following it, is possible to integrate and start nonnative systems in QEMU. Docker is advantageous to setup and start systems fast. It would be nice to integrate the docker file system into QEMU. After building a docker image, it is realizable to export the file system into a directory with the name **rootfs** with the command

```
docker export $(docker create initrd-s390x) | tar -C "rootfs"-xvf -.
```

Linux has got the feature that it is adventitious to reformat directories for file systems and to copy/ mount content into this one. Reformatting the default docker file system UnionFS to another one as ext4 for example can be done then.

QEMU is accepting the new file system as a block device for the guest system then. The default path to the mounted file system as a hard disk is `/dev/vda` as the first partition [Whit 20, p.22].

### 2.5.1. UnionFS

Docker does not use any default Linux file system. Docker images are based on the Union File System (UnionFS) [Ashr 15, p.21]. This file system is using different file system layers with grouping directories and files in branches. The first layer is the typical Linux boot file system with the name `bootfs`. That performs the same as in a Linux virtualization stack with using memory at first and unmounting to receiving RAM free by the `initrd` disk image. So the `bootfs` can be used inside of another Linux file system to mount in an virtualization stack for a successful boot process with QEMU. The next layer contains the base image with the operating system given by the "FROM" command. Then every docker command inside of

the Dockerfile is adding an additional layer with the installation of applications or building binary files. That is the reason that every executed docker command is receiving an own id in the disk memory during the build process. Every separate docker command is using his own disk space. So a docker image can grow really fast. It is reasonable to compress so much as possible of different routines into one docker command. All sizes of different docker layers will exist continuously inside of the new file system.

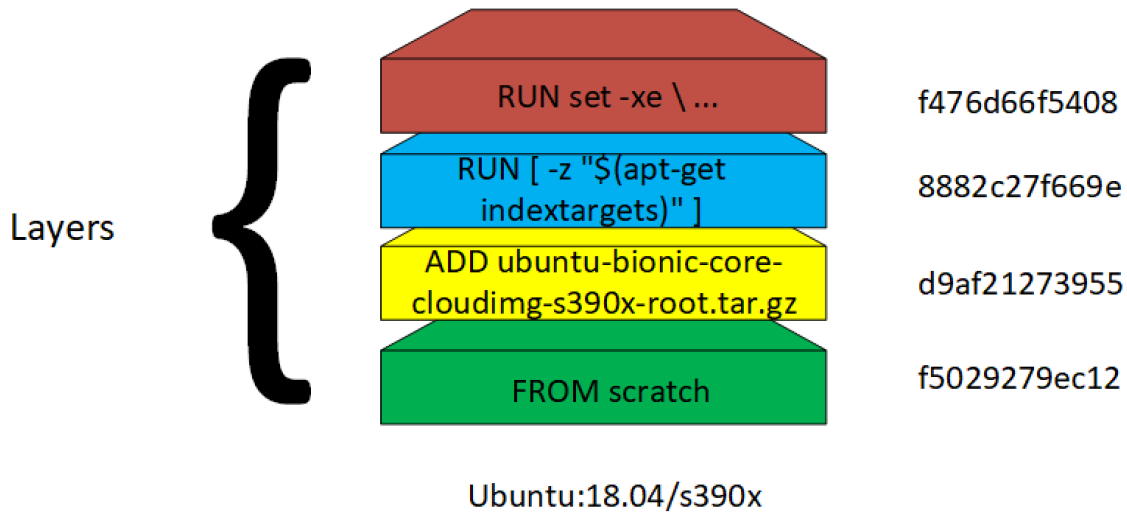


Figure 2.3.: Union File System

### 2.5.2. Ext4

Ext4 is a journaling file system (the same as ext3). The journal is registering transactional all changes on the operating system with meta data. So no data are lost after a system failure. They can be restored based on the journal without a save procedure by the user. Such journal file systems in the ext file system family are working with blocks [Seuf15, p.20]. The journal is splitted into the journal super block, descriptor blocks, commit blocks and revoke blocks. The super block contains all meta data of the journal. Descriptor blocks include the special destination adress and a sequence number. Flex groups can combine multiple group blocks - inode bitmaps, inode tables and data blocks - to one logic block group. All data from inodes are only saved in the first flex group then. So the search for files is more powerful because meta data are allocated at one place.

All meta data from the journal can be relocated in inodes because changes of file system meta data are registered, too.

The difference to the journal in ext3 are the option writing asynchronous commit blocks and additional check sums for journal transactions [Seuf 15, p.28].

Ext4 provides a better performance than ext3. This file system can be formatted with `mkfs.ext4` which is available in every Linux distribution. This tool is creating a group desriptor table for further group descriptors what allows an expanding of the file system. The file system can grow a maximum of the multiple 1024 of his existing size because of the saved space [Seuf 15, p.21].

Another feature of ext4 is the possibility of inline files and inline directories. So small files and directories can be saved directly in inodes instead of data blocks. From this follows less disk space consumption [Seuf 15, p.24].

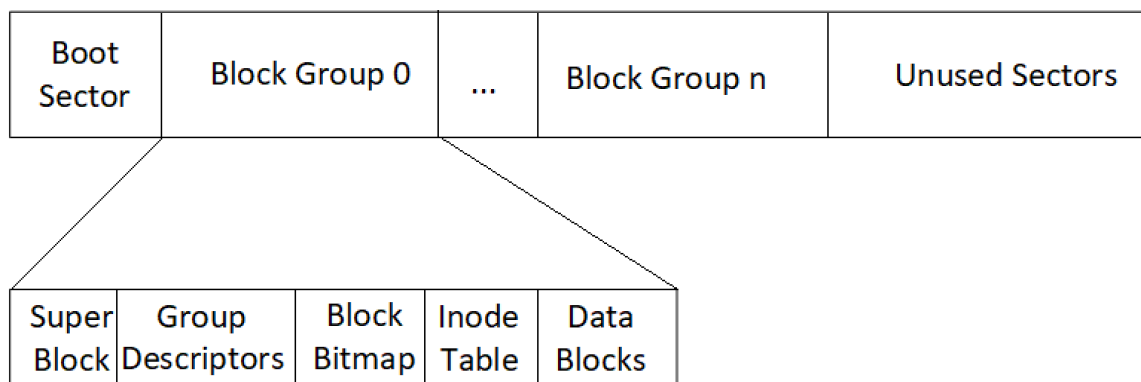


Figure 2.4.: Ext4

### 2.5.3. Mounting an External File System

In 2.5 "File Systems" is described how to export a Docker file system into a directory. That has got the format of the 2.3 "UnionFS" at this moment. QEMU contains the beneficial tool **qemu-img** for creating loadable images of file systems. These require the minimal size of the docker image. The output of the size is given by `docker images | grep ${id}`. Qemu-img understands only rounded numbers. Therefore, that has to be rounded up to a full number. That will be done automated with the tool **awk**. In the case of Cassandra, there will be used `docker images | grep 'Cassandra:s390x' | awk '{print int($7+0.5)"G"}'`. `print $7` would give the output of the size raw of `docker images`, which is a float number. The `int` is rounding down to an integer number. Accordingly, the `0.5` has to be summed up for rounding up. The GB has been removed in this process. **qemu-img** requires a G behind the number. That can be added with "G" inside of the awk command.

Finally, Cassandra is using the command

```
qemu-img create -f raw cassandra.img 2G
```

 (see 6.6 "Running the Docker Image")

for the creation of a filesystem image with a Docker image size about 1.2G then. "2G" can be replaced by the command above in the CI/CD pipeline then (see A.1 "Circle CI configuration for Apache Cassandra").

At first, `cassandra.img` is empty. The content of **rootfs** has to be transferred into this image file. But QEMU does not know UnionFS as a file system. That is the reason for formatting `cassandra.img` with `mkfs.ext4 -F cassandra.img` to ext4. Writing into `cassandra.img` is only executable if the image is mounted into a directory under `/mnt/` with

```
mount -o loop cassandra.img /mnt/rootfs.
```

In the next step the content of the file system directory can be copied into `/mnt/rootfs/` with

```
cp -r rootfs/* /mnt/rootfs/..
```

Then all data of the Docker image are saved into `cassandra.img` because it is mounted into `/mnt/rootfs/`. In conclusion, the image `cassandra.img` is usable for deployments with **qemu-system-s390x** (see 6.5 Run Cassandra).

## Chapter 3.

### Prerequisites

Different software is necessary to run qemu or docker for multiple architectures. Therefore, **docker** and **qemu** should be installed. Additionally, `qemu-user-static` <sup>1</sup> and `binfmt_misc` <sup>2</sup> are important for running multi-architecture containers.

**Root permissions** are required for installation, configuration and running emulated systems.

It is possible to install packages as **binfmt-support** and **qemu-user-static** by different Linux distributions, but it is recommended to use the latest version for s390x.

#### 3.1. Registration of Qemu-S390x

The packages **qemu-user-static** and **binfmt-support** should be installed because of the configuration.

The Linux kernel module `binfmt_misc` can be mounted with the following command:

```
mount binfmt_misc -t binfmt_misc /proc/sys/fs/binfmt_misc
```

The installed version 3.1 of `qemu-user-static` on Ubuntu 18.04 contains bugs that Docker images for s390 are not buildable on x86. Therefore, an upgrade to a new release is necessary. Ultimate stable releases of `qemu-user-static` can be found under <https://github.com/multiarch/qemu-user-static/releases/>. The release v5.0.0-2 is used for the project and downloaded with

```
wget https://github.com/multiarch/qemu-user-static/releases/download/v5.0.0-2/x86_64_qemu-s390x-static.tar.gz
```

for the specific version of `qemu-s390x-static` on x86. That is extracted to the directory `/usr/bin/` with the command

```
sudo tar -xvzf x86_64_qemu-s390x-static.tar.gz -C /usr/bin/
```

---

<sup>1</sup><https://github.com/multiarch/qemu-user-static>

<sup>2</sup><https://www.kernel.org/doc/html/latest/admin-guide/binfmt-misc.html>





```
git clone git://github.com/docker/buildx && cd buildx
make install
```

Figure 3.3.: BuildX Upgrade

### 3.3. Fetching Signed Linux Kernel Image

QEMU needs a built Linux kernel to start a full system emulation. There are the two options, to build the Linux kernel repeating because of security updates or to fetch the matching Linux kernel for the Docker image from the associated repository of the corresponding Linux distribution.

The build process via CI/CD should be fast. Therefore, the second option will be used. The Docker filesystem in rootfs does not contain any Linux kernel in the directory `/boot/`. The prerequisite for fetching of a signed Linux kernel image for a minimal system is the installed package **debootstrap**. This command line tool can download the filesystem of a minimal system to a sub directory on a Debian based system. There are additional options to specify the architecture with `--arch=s390x` in our case and to include the Linux kernel with `--include=linux-generic`. To receive only a minimal system, the additional option `--variant=minbase` should be used. This command "debootstrap" is validating all downloaded packages additionally. This process is easy to automate and requires only 10 seconds. In conclusion, the generic and signed Linux kernel image can be used in the `/boot/` directory for the integration into QEMU.

### 3.4. Optimized QEMU Command

Every additional device requires additional performance and time for starting the system. So the systems requirements had to be figured out to be minimal for every given open-source project and for running tests on it. That counts for the number of CPUs, too.

The kernel option is receiving the path to the built s390x kernel with the name bzImage. In the given 3.4 "Optimized QEMU command" the qemu command is executed in the directory with the built Linux kernel and the path is unimportant.

The option **-m** is available to add the minimal guest memory matching the system requirements of every open-source project. **-M s390-ccw-virtio** defines the machine type for s390x. **-nodefaults** is deactivating default additional devices activated in QEMU. Only the console is necessary for receiving an output and debugging. So this one is added as a device with **-chardev stdio,id=console,signal=off,mux=on -mon chardev=console** and **-device sclpconsole,chardev=console**. Additionally, the console has to be specified

with **console=ttyS0** in the append part.

Cassandra as a project does not need any network interface (-net none) or parallelism (-parallel none). The option **-nographic** is responsible for not adding any graphical interface. So we save system requirements. The option **-smp** is the minimal number of CPUs for the guest. The file system of containers can be loaded as a hard disk with the option **-hda** which is explained in every chapter of the Bachelor Thesis for a considered open-source project. That is the ideal option to mount a minimal file system for every application or system. **/dev/vda** is the partition name and **rdinit** is used for using bash as a default shell. This command can be used for user mode emulation with Apache Cassandra as an example.

```
/usr/bin/qemu-system-s390x -kernel bzImage -m 4G -M s390-ccw-virtio -nodefaults \  
-device sclpconsole,chardev=console -parallel none -net none -chardev stdio,\  
id=console,signal=off,mux=on -mon chardev=console -nographic -smp 3 \  
-hda /data/cassandra.img --append 'root=/dev/vda rw console=ttyS0 \  
rdinit=/bin/bash'
```

Figure 3.4.: Optimized QEMU Command

## **Chapter 4.**

# **Continuous Integration**

### **4.1. CI/CD in the Software Development Lifecycle**

CI/CD is the short name for Continuous Integration/Continuous Delivery. This method contains systems and technologies used for Test Driven Development. Continuous Delivery provides executing changes in the code, building them to an application and testing it directly. If all is working fine, it can be merged. Continuous Delivery is adding automated releases to the process. The software is tested automatically by the system before. If all is ok, it can be released.

These tests during the development are integrated into all sprints for agile software development. That has got the name DevOps after adapting these practices with expanding with automated deployments for operations. It is playing an important role through all phases in the software development because of the quality of code today. The code complexity is growing with the size of the project. Therefore, it is beneficially to have automated tests running directly after the commit.

### **4.2. CI/CD Systems**



## Chapter 5.

# Kubernetes

### 5.1. Overview

Kubernetes is a container platform for high availability clusters. There exist plugins for integration tests for Kubernetes with the name kubetest<sup>1</sup>. They contain conformance tests, as well as e2e tests (end-to-end), too. That can be all built and executed on the system. Therefore, a Dockerfile for setting up Kubernetes and building tests with Go is necessary. The challenge is, that 2 big Github repositories have to be cloned and integrated into the docker image. That is using a lot of space. The solution is using a multi staging Dockerfile. So 2 different Dockerfiles are used in one Dockerfile and one of them is used for building. The other one is used for the installation and testing with built tests. At the end the size of the docker image has got only the size of the test image regardless of the repository size in the mother Dockerfile.

### 5.2. Multi-Staging Dockerfile

A multi-staging Dockerfile is using different systems in one Dockerfile for different stages. These systems are receiving special names as indicators with "AS" behind the "FROM" with the base image name. Default this feature is used for building applications in one stage and executing the copied application in another stage. The same counts for cloning Github repositories and building binary files based on it. On this way, a lot of space is saved. Concluding, the docker image has got only the size of the executing system with the application file (without all the code). That is an "experimental feature" at the moment. Therefore the **experimental flag** is necessary to export or set before using it (see 3.2 Multi-Architecture Images).

Default one image is receiving the name build and the other one a name of what will be done with that. In our case, the second image has got the name work. The second image is copying with "COPY --from=build" all required built data from the first image that it can

---

<sup>1</sup><https://kubetest.readthedocs.io/en/latest/>

be used for running the application. On this way, it is possible to reduce the size of a docker image.

Multi-staging Dockerfiles are an approved method for executing binaries based on Github repositories.

### 5.3. Installation

Kubernetes needs a lot of packages for running and for tests. That will be all installed with the RUN command. [apt.kubernetes.io](https://apt.kubernetes.io) has got later versions of Kubernetes than the Ubuntu repository. Therefore this repository has to be added to Ubuntu. kub-build is the name of the mother Dockerfile to be able to copy needed files and directories from there.

In the intallation part of the Dockerfile the Kubernetes repository <https://apt.kubernetes.io> for Debian packages has to be registered together with with the gpg key used by <https://packages.cloud.google.com/apt/doc/apt-key.gpg>. A Dockerfile is installing only necessary packages. Therefore, apt-transport-https, apt-utils, curl, git, ca-certificates, gnupg-agent and software-properties-common have to be installed first for the following installation. The system requires the latest update with `apt-get update` after the registration of the additional repository besides of the default Ubuntu repositories imported with the base Ubuntu image "s390x/ubuntu:18.04" in the FROM command. After that the packages docker.io, kubelet and kubeadm can be installed from kubernetes.io. **Docker.io** contains the container engine docker with all docker commands. CRI/O or containerd would be allowed, too. The Docker daemon will be used because that is the main used container engine of the Kubernetes project and all tests are running withit.

**Kubelet** is the primary node agent running on each node. He is responsible that different containers can run together in a pod. Pods are deployable units defined in JSON or a yaml file. They include one or a group of containers with shared storage and network resources. Hosting of distributed systems with different services in different containers can work together. Consequential one pod is something as one "logical host".

**Kubeadm** is the administration tool to setup clusters. It is necessary to upgrade Kubernetes to other versions, too. Clusters can be initialized. The network can be configured and the command **kubectl** (Kubernetes Control Plane) for adding additional nodes to a cluster can be initialized.

`apt-mark hold` is keeping these special versions of kubelet, kubeadm and kubectl.

```

FROM s390x/ubuntu:18.04 AS kub-build

# The author
MAINTAINER Sarah Julia Kriesch <sarah.kriesch@ibm.com>

#Installation
RUN echo "Installing necessary packages" && \
apt-get update && apt-get install -y \
apt-transport-https \
apt-utils \
systemd \
curl \
git \
ca-certificates \
gnupg-agent \
software-properties-common \
&& curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add - \
&& echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" \
> /etc/apt/sources.list.d/kubernetes.list \
&& apt-get update && apt-get install -y \
docker.io \
kubernetes \
kubeadm \
&& apt-mark hold kubernetes \
&& apt-get clean \
&& rm -rf /var/lib/apt/lists/* /tmp/* /var/tmp/* \
&& systemctl enable docker

```

Figure 5.1.: Kubernetes Installation

## 5.4. Installation of the Ultimate Go

There were some issues with older Go versions as 1.10 during building tests for Kubernetes. Therefore a higher version (min. 1.13) should be used. It is recommended to use the ultimate go version for last versioned Kubernetes tests. It is possible to receive the version number of the ultimate go release with the command

```
curl https://golang.org/VERSION?m=text.
```

This version number has to be included before linux-s390x.tar.gz for downloading the special s390x archive from the go directory by [dl.google.com](https://dl.google.com). Then the version number has to

be called with curl inside of another curl command with the whole path to the special tar archive on [dl.google.com](https://dl.google.com). Every tar archive has got the same structure for every version (`$version.$architecture.tar.gz`). On this way the latest version of Go is integrable into the curl command that it can be installed. Directories for bin, pkg and src have to be created after extracting this tar archive in the `/root/` directory. They are not integrated in the tar archive.

The environment variables for GOROOT, GOPATH and PATH have to be set with ENV on the top of the Dockerfile for successful builds later. PWD is added because Github repositories have to be cloned to this directory.

The ENV variables will be on the top of the Dockerfile. The part for the "Installation of Go" will be attached to the end of the Kubernetes installation part.

```
ENV GOROOT=/root/go
ENV GOPATH=/root/go
ENV PATH=$GOPATH/bin:$PATH
ENV PATH=$PATH:$GOROOT/bin
ENV PWD=/root/go/src/

#Installation of latest GO
&& echo "Installation of latest GO" && \
curl "https://dl.google.com/go/ \
$(curl https://golang.org/VERSION?m=text).linux-s390x.tar.gz" \
| tar -C /root/ -xz \
&& mkdir -p /root/go/{bin,pkg,src}
```

Figure 5.2.: Go Installation

## 5.5. Building Tests

After a successful installation of go, it is possible to build and install the Kubernetes test environment. At first the directory `k8s.io` has to be created because `kubernetes-tests` are looking for this directory as a mother directory. The repository `test-infra` by the Kubernetes project has to be cloned to there. The repository **test-infra** contains all tests for Kubernetes provided by the community. These can be used of course and are updated continuously. That is the reason to clone this repository inside of the Dockerfile. Inside of this `test-infra` directory `kubetest` can be installed with **go install**. That is downloading all available Kubernetes-Tests.



So you can use them to test the own Kubernetes cluster and the used software.

The name of the most relevant tests for the Kubernetes community is "conformance tests". These conformance tests are executed with `e2e.test` which can be built with `make` inside of the `kubernetes` repository. Therefore this repository has to be cloned to `k8s.io`, too. These tests certify the software to comply regular standards. Only with complying these standards, Kubernetes software is allowed to become Kubernetes certified<sup>2</sup>.

```
&& cd $PWD \  
#Clone test-infra  
&& mkdir -p $GOPATH/src/k8s.io \  
&& cd $GOPATH/src/k8s.io \  
&& git clone https://github.com/kubernetes/test-infra.git \  
/root/go/src/k8s.io/test-infra \  
&& cd /root/go/src/k8s.io/test-infra/ \  
#Install kubetest  
&& GO111MODULE=on go install ./kubetest \  
#Build test binary  
&& git clone https://github.com/kubernetes/kubernetes.git \  
/root/go/src/k8s.io/kubernetes \  
&& cd /root/go/src/k8s.io/kubernetes/  
  
CMD make WHAT="test/e2e/e2e.test vendor/github.com/onsi/ginkgo/ginkgo cmd/kubect1"
```

Figure 5.3.: Test Building for Kubernetes

### 5.5.1. E2e.test

## 5.6. Run in the Main Dockerfile

## 5.7. Integration into CI/CD

---

<sup>2</sup><https://github.com/cncf/k8s-conformance>



## Chapter 6.

# Cassandra

### 6.1. Overview

Apache Cassandra is a Java based NoSQL database management system. It contains the commandline interface CQL for Cassandra database commands instead of SQL. It has to be recognized that the supported Java version by the open-source project Apache Cassandra can be changed with every version. That has to be recognized during the integration into the CI/CD pipeline of the project. There are different multi-arch Dockerfiles for all versions of AdoptOpenJDK to use.

This Apache project is using Circle CI for Continuous Integration.

### 6.2. Installation

It should be created a Dockerfile<sup>1</sup> for a multi-arch image in this project with the reason for using it for different architectures. Apache Cassandra is a Java based application. Therefore, the research has contained the search for a multi-arch image with pre-installed Java. The Cassandra version 3.11.6 is supporting Java 8. AdoptOpenJDK is maintaining multi-arch images for different Linux distributions with pre-installed AdoptOpenJDK. In the description of [2.2](#)"Multi-Arch Image AdoptOpenJDK 8" is explained how the recommended Java version will be chosen for the specific architecture in this base image. This base image is integrated with `FROM adoptopenjdk:8-jdk-hotspot` on top of the Dockerfile.

Apache Cassandra will be cloned from the Github repository with  
`git clone https://github.com/apache/cassandra.git`.

The special version is set with an agument with the ARG variable. On this way, it is possible to checkout other versions with a given argument as `--build-arg CASSANDRA_VERSION=3.12.4` in the docker build command.

Apache Cassandra is using JNA (Java Native Library)<sup>2</sup> as an extension of Java. JNA offers

---

<sup>1</sup><https://github.com/s390x-container-samples/s390x-cassandra-ci-cd/blob/master/Dockerfile>

<sup>2</sup><https://github.com/java-native-access/jna>

provides Java programs easy access to native shared libraries. The ant command is building the application and the JNA subsequently.

The command `sed -i ' s/Xss256k/Xss32m/'` is setting the size of the frame stack (JVM stack) used by each thread to store local variables. This frame size is reduced to a minimum and stored in both files `build.xml` and `conf/jvm.options`.

```
FROM adoptopenjdk:8-jdk-hotspot
ARG CASSANDRA_VERSION=3.11.6
git clone https://github.com/java-native-access/jna.git \
&& cd jna \
&& git checkout 4.2.2 \
&& ant native jar \
# Build and install Apache Cassandra
&& cd $SOURCE_ROOT \
&& git clone https://github.com/apache/cassandra.git \
&& cd cassandra \
&& git checkout cassandra-${CASSANDRA_VERSION} \
&& sed -i ' s/Xss256k/Xss32m/' build.xml conf/jvm.options \
&& ant \
&& rm lib/snappy-java-1.1.1.7.jar \
&& wget -O lib/snappy-java-1.1.2.6.jar https://repo1.maven.org/maven2/org/xerial/snappy/
snappy-java/1.1.2.6/snappy-java-1.1.2.6.jar \
&& rm lib/jna-4.2.2.jar \
&& cd $SOURCE_ROOT \
&& cp $SOURCE_ROOT/jna/build/jna.jar $SOURCE_ROOT/cassandra/lib/jna-4.2.2.jar \
&& cp -R $SOURCE_ROOT/cassandra /usr/local/ \
```

Figure 6.1.: Cassandra Installation

### 6.3. Java Optimization

Apache Cassandra should be allowed to use all available memory of the minimal system effectively. Default the JVM is configured that only a part of that can be used. These configurations are set in the Java Options. Afterwards, Apache Cassandra can work with a good performance.

Every major Java release can differ from another one with such options for Java optimization.

Following Java optimizations are available for Java 8:

- **-XX:+UseCGroupMemoryLimitForHeap**

This option tells the JVM to look into `/sys/fs/cgroup/memory/memory.limit` after the available memory and to use the whole memory if necessary [Flo0 17].

- **-XX:ParallelGCThreads=2**

Apache Cassandra is referencing 2 required CPUs as a minimum in their documentation [Cass]. This JVM option permits 2 GC (Garbage Collection) threads parallel on different (virtual) processors.

All these Java optimization options can be set as an environment variable with `ENV JVM_OPTIONS` in the Dockerfile. Then they will be used for every Java process like Apache Cassandra as an example. As a result, Apache Cassandra can use the whole memory and delivered CPU, which is specified in the QEMU options.

## 6.4. Nagios Monitoring Check for Tests

Nagios is an open-source monitoring system to check the status of systems with their services. There are many checks for different services available. One monitoring check<sup>3</sup> exists for Apache Cassandra. It is testing whether Java would be working, the service Apache Cassandra would be up and running, and whether CQL commands would be possible to execute without any issues. This Perl script is using **nodetool** for monitoring the JVM and the application. Nodetool is delivered together with Cassandra. It is connecting with the database and can dispense statistics about the Cassandra cluster / host [Carp 20, p.256]. The result contains the status, information about memory usage and other capacities. This tool can indicate issues resulting in messages via the monitoring plugin in our case.

This monitoring check is written in Perl and is a nice choice for automated tests with error messages in the case if any output would not give an **OK**.

All monitoring plugins in the referenced Github repository are developed additionally for integration into Docker. Following, these scripts do not require any installed nagios system for monitoring. Summarily, these monitoring plugins can be used for tests inside of our Dockerfile.

That said, **cassandra\_check.pl** is integrated into the Dockerfile with following commands at the end:

```
COPY plugins/cassandra_check.pl /bin/cassandra_check.pl
ENTRYPOINT ["perl", "cassandra_check.pl"]
```

Figure 6.2.: Cassandra Monitoring Check

<sup>3</sup>[https://github.com/skyscrapers/monitoring-plugins/blob/master/check\\_cassandra.pl](https://github.com/skyscrapers/monitoring-plugins/blob/master/check_cassandra.pl)

## 6.5. Workaround Because of an JVM Issue

Apache Cassandra is using OpenJDK 8 for running as default. There are some discussions<sup>4</sup> about supporting OpenJDK 11 at the moment. The Dockerfile mentioned above is using AdoptOpenJDK 8 at the build process is working without any problems on Z systems. There exists a JVM issue during the build process on x86 which is the reason for non successful builds of Java applications there. That counts for all more complex Java applications.

This case has been tested with AdoptOpenJDK 8 and AdoptOpenJDK 11. The library `hsdis-s390.so` is not loadable for both Java versions on x86, but during the build process on s390x. This issue does not exist with the version AdoptOpenJDK 14, but Apache Cassandra does not work with this version.

That is an emulation bug of qemu which has to be fixed. This issue has been reported to KVM Developers inside of IBM.

---

<sup>4</sup><https://lists.apache.org/thread.html/r38f6beaa22e247cb38212f476f5f79efdc6587f83af0397406c06d7c%40dev.cassandra.apache.org>

```

#
# A fatal error has been detected by the Java Runtime Environment:
#
# SIGILL (0x4) at pc=0x0000004012539240, pid=18, tid=21
#
# JRE version: OpenJDK Runtime Environment (11.0.7+10) (build 11.0.7+10)
# Java VM: OpenJDK 64-Bit Server VM (11.0.7+10, mixed mode, tiered, compressed oops,
g1 gc, linux-s390x)
# Problematic frame:
# J 1 c1 java.lang.Object.<init>()V java.base@11.0.7 (1 bytes) @ 0x0000004012539240
[0x0000004012539200+0x0000000000000040]
#
# Core dump will be written. Default location: //core
#
# An error report file with more information is saved as:
# //hs_err_pid18.log
Compiled method (c1)      943      1      3      java.lang.Object::<init> (1 bytes)
  total in heap [0x0000004012539010,0x00000040125393b8] = 936
  relocation    [0x0000004012539178,0x00000040125391a8] = 48
  constants     [0x00000040125391c0,0x0000004012539200] = 64
  main code     [0x0000004012539200,0x0000004012539300] = 256
  stub code     [0x0000004012539300,0x0000004012539358] = 88
  metadata      [0x0000004012539358,0x0000004012539370] = 24
  scopes data   [0x0000004012539370,0x0000004012539380] = 16
  scopes pcs    [0x0000004012539380,0x00000040125393b0] = 48
  dependencies  [0x00000040125393b0,0x00000040125393b8] = 8
Compiled method (c1)      948      1      3      java.lang.Object::<init> (1 bytes)
  total in heap [0x0000004012539010,0x00000040125393b8] = 936
  relocation    [0x0000004012539178,0x00000040125391a8] = 48
  constants     [0x00000040125391c0,0x0000004012539200] = 64
  main code     [0x0000004012539200,0x0000004012539300] = 256
  stub code     [0x0000004012539300,0x0000004012539358] = 88
  metadata      [0x0000004012539358,0x0000004012539370] = 24
  scopes data   [0x0000004012539370,0x0000004012539380] = 16
  scopes pcs    [0x0000004012539380,0x00000040125393b0] = 48
  dependencies  [0x00000040125393b0,0x00000040125393b8] = 8
Could not load hsdis-s390x.so; library not loadable; PrintAssembly is disabled
#
# If you would like to submit a bug report, please visit:
#   https://github.com/AdoptOpenJDK/openjdk-support/issues
#
Aborted (core dumped)

```

Figure 6.3.: JVM Build Issue

There is a workaround for the development process for finishing the project. The Docker image is buildable on s390x. Therefore, the Docker image has been built for s390x on s390x with

```
docker build --squash -t cassandra ..
```

Afterwards, this image can be saved as a tar archive with

```
docker save cassandra > cassandra.tar
```

on the host. It has to be transferred with rsync to the x86 system then. The image inside of the tar archive can be included to all existing images of `docker images` with the command `docker load --input cassandra.tar`.

After this action, the Docker image is usable as it would have been built on this system. It can be integrated into QEMU and can be stated without any problems afterwards.

## 6.6. Run Docker Image in QEMU

The Docker image will be built with the command

```
docker build buildx --platform=linux/s390x --squash -t cassandra:s390x .
```

in the directory with the Cassandra Dockerfile for the architecture s390x.

**Squash** is an option to compress a Docker image and combine commands in a Dockerfile automatically. The prerequisites for building s390x images on x86 are set during the emulation preparation. The command `docker images` has to show the registered Docker image with the name `cassandra:s390x` then. It should be possible to integrate this Docker image into the `qemu` command. Therefore, a `qemu-image` will be created with an rounded given size besides of the Docker image in the `docker images` command. So the command

```
qemu-img create -f raw cassandra-s390x.img 2G
```

can be used. This image needs any Linux file system because QEMU does not know the Docker file system. The image is formatted with the command `mkfs.ext4 -F cassandra-s390x.img` then. For receiving the file system of the docker image a directory with the name `rootfs` has to be created and the command

```
docker export $(docker create cassandra:s390x) | tar -C "rootfs"-xvf -
```

is exporting the docker image into the directory `rootfs`. Following transfers the content of `rootfs` into the image `cassandra.img`.

```
mkdir /mnt/rootfs
mount -o loop cassandra-s390x.img /mnt/rootfs
cp -r rootfs/* /mnt/rootfs/.
```

Figure 6.4.: Mount Rootfs



Now it is possible to run the system with Cassandra:

```
/usr/bin/qemu-system-s390x -kernel bzImage -m 40G -M s390-ccw-virtio -nodefaults \
-device sclpconsole,chardev=console -parallel none -net none -chardev stdio, \
id=console,signal=off,mux=on -mon chardev=console -nographic -smp 3 \
-hda /data/dockerfile-examples/ApacheCassandra/cassandra.img \
--append 'root=/dev/vda rw console=ttyS0 rdinit=/bin/bash'
```

Figure 6.5.: Run Cassandra

## 6.7. Integration into CI/CD

The community of Apache Cassandra is using Circle CI<sup>5</sup> as a Continuous Integration system. This CI/CD system is connected with Github. The tests are running after every commit in the master repository. These tests can run in a container or in a virtual machine. These options make it possible to build the Docker image for the architecture s390x and to integrate it into QEMU via the CI/CD pipeline without additional scripting or workarounds.

Circle CI is working with a yaml configuration<sup>6</sup>. An enhancement is the opportunity to define shell commands with "command" in the yaml file. Therefore, "3.3 Fetching Signed Linux Kernel Image" can be integrated. Additionally, dependencies can be created. On this way, it can be assured that the build process for the Docker image is triggered before the filesystem creation. The same counts for the kernel fetching process and the filesystem creation before the QEMU start.

This configuration file with the whole test process has got the name **config.yml**.

**Workflows** defines the order of different process steps with their dependencies with **-requires**. The different processes are specified separately above then. The workspace can be declared. Running steps can receive names. Shell commands can be executed. Additionally, binaries and other files can be stored as artifacts.

The implementation of the job configuration can be found under [A.1](#).

<sup>5</sup><https://circleci.com/>

<sup>6</sup><https://circleci.com/docs/2.0/sample-config/>



## Chapter 7.

### Outlook

This project has been combining container technologies with an emulator which is used as a foundation for different virtualization technologies. Additionally, it has represented a method to deploy applications fast with an emulated system for alternative system architectures. Emulation bugs for QEMU have to be fixed. Afterwards, this method can be used in all CI/CD pipelines of open-source projects and on computers by Developers. New features can be tested without high effort. Following, the software development process can become faster. There is the possibility to embed the QEMU deployment of Docker images into virtualization technologies as KVM and XEN. But that requires additional development for alignments. All in all, this emulation method should be usable for fast deployments in CI/CD pipelines and system deployments for all system architectures.



## Chapter 8.

### Summary

Hier kommt eine Zusammenfassung





## Appendix A.

### Supplemental Information

#### A.1. Circle CI configuration for Apache Cassandra

```

jobs:
  build-docker-image:
    working_directory: ~/s390x
    docker:
      - image: Dockerfile
    steps:
      - attach_workspace:
          at: .

      - run:
          name: Build for s390x
          command: docker build buildx --platform=linux/s390x --squash \
            -t cassandra:s390x .

      - docker/check:
          registry: $DOCKER_REGISTRY
  prepare-kernel:
    steps:
      - run:
          name: Fetching the Linux kernel
          command: |
            mkdir ~/s390x/kernel-bionic; debootstrap --include=linux-generic \
              --arch=s390x --variant=minbase bionic kernel-bionic \
              http://ports.ubuntu.com

      - store_artifacts:
          path: ~/s390x/kernel-bionic/boot/vmlinuz-4.15.0-20-generic

      - persist_to_workspace:
          root: .
          paths:
            - .

```



```

prepare-filessystem:
  working_directory: ~/s390x
  steps:
    - attach_workspace:
        at: .
    - run:
        name: Creating a ext4 filesystem with integration of the Cassandra image
        command: |
            mkdir rootfs; qemu-img create -f raw cassandra-s390x.img \
            $(docker images | grep 'cassandra:s390x' | \
            awk '{print int($7+0.5)"G"}'); \
            docker export $(docker create cassandra:s390x) | \
            tar -C "rootfs"-xvf -; mkfs.ext4 -F cassandra-s390x.img;\
            mount -o loop cassandra-s390x.img /mnt/rootfs; \
            cp -r rootfs/* /mnt/rootfs/.
    - store_artifacts:
        path: ~/s390x/cassandra-s390x.img
test:
  working_directory: ~/s390x
  steps:
    - attach_workspace:
        at: .
    - run:
        name: Run test with QEMU
        command: |
            /usr/bin/qemu-system-s390x -kernel vmlinuz-4.15.0-20-generic \
            -m 4G -M s390-ccw-virtio -nodefaults -device sclpconsole,\
            chardev=console -parallel none -net none -chardev stdio,\
            id=console,signal=off,mux=on -mon chardev=console \
            -nographic -smp 3 -hda cassandra-s390x.img \
            --append 'root=/dev/vda rw console=ttyS0 rdinit=/bin/bash'
workflows:
  version: 1
  build-test:
    jobs:
      - prepare-kernel
      - build-docker-image
        requires:
          - prepare-kernel
      - prepare-filessystem
        requires:
          - build-docker-image
      - test
        requires:
          - prepare-kernel
          - prepare-filessystem

```



## List of Figures

2.1. Emulation with QEMU . . . . .	9
2.2. Multi-Arch Image AdoptOpenJDK 8 . . . . .	12
2.3. Union File System . . . . .	14
2.4. Ext4 . . . . .	15
3.1. Register S390x Binaries . . . . .	18
3.2. Docker Experimental Flag . . . . .	18
3.3. BuildX Upgrade . . . . .	19
3.4. Optimized QEMU Command . . . . .	20
5.1. Kubernetes Installation . . . . .	25
5.2. Go Installation . . . . .	26
5.3. Test Building for Kubernestes . . . . .	27
6.1. Cassandra Installation . . . . .	30
6.2. Cassandra Monitoring Check . . . . .	31
6.3. JVM Build Issue . . . . .	33
6.4. Mount Rootfs . . . . .	34
6.5. Run Cassandra . . . . .	35



## List of Tables



## Bibliography

- [Ashr 15] W. Ashraf. *The Docker EcoSystem*. Gitbook, 2015.
- [Barb 18] D. H. Barboza. “Enhancing QEMU virtio-scsi with Block Limits vital product data (VPD) emulation”. October 2018. <https://developer.ibm.com/technologies/linux/articles/enhancing-qemu-virtio-scsi-with-block-limits-vpd-emulation/#sec3.1> [Accessed 17 August 2020].
- [Bloc 19] B. Block, Ed. *Modern Mainframes & Linux Running on Them*, IBM Deutschland Research & Development GmbH, Chemnitzer Linux-Tage, March 2019. <https://chemnitzer.linux-tage.de/2019/media/programm/folien/173.pdf> [Accessed 07 February 2020].
- [Butt 11] K. Butt, A. Qadeer, and A. Waheed. “MIPS64 User Mode Emulation: A Case Study in Open Source Software Engineering”. In: *2011 7th International Conference on Emerging Technologies*, pp. 1–6, 2011.
- [Carp 20] J. Carpenter and E. Hewitt. *Cassandra: The Definitive Guide*. "O'Reilly Media, Inc.", 2020.
- [Cass] A. Cassandra. “Apache Cassandra Documentation: Hardware Choices”. <https://cassandra.apache.org/doc/latest/operating/hardware.html> [Accessed 10 September 2020].
- [Cota 17] E. Cota, P. Bonzini, A. Béné, and L. Carloni. “Cross-ISA machine emulation for multicores”. In: *2017 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, pp. 210–220, February 2017.
- [Data] Datastax. “What is Cassandra?”. <https://www.datastax.com/cassandra> [Accessed 26 August 2020].
- [Floo 17] C. Flood. “OpenJDK and Containers”. April 2017. <https://developers.redhat.com/blog/2017/04/04/openjdk-and-containers/#more-433899> [Accessed 10 September 2020].

- [Gün] R. Günther. “Kernel Support for miscellaneous Binary Formats (binfmt\_misc)”. <https://www.kernel.org/doc/Documentation/admin-guide/binfmt-misc.rst> [Accessed 22 August 2020].
- [Lasc 20] O. Lascu, B. White, J. Troy, and F. Packheiser. *IBM z15 Technical Instruction*. IBM Redbooks, 2020.
- [Linu] O. M. P. a Linux Foundation Project. “The Modern Mainframe”. <https://www.openmainframeproject.org/modern-mainframe> [Accessed 04 September 2020].
- [Opsa 13] J. M. Opsahl. *Open-source virtualization - Functionality and performance of Qemu/KVM, Xen, Libvirt and VirtualBox*. Master’s thesis, University of Oslo, May 2013.
- [QEMU] QEMU. “QEMU User space emulator”. <https://www.qemu.org/docs/master/user/main.html> [Accessed 10 September 2020].
- [Rose 15] D. S. Rosenthal. *Emulation & Virtualization as Preservation Strategies*. Andrew W. Mellon Foundation, 2015.
- [Scho 19] B. Scholl, T. Swanson, and P. Jausovec. *Cloud Native: Using Containers, Functions, and Data to Build Next-Generation Applications*. "O’Reilly Media, Inc.", 2019.
- [Seuf 15] S. Seufert. *Implementierung eines forensischen Ext4-Incode-Carvers*. Master’s thesis, Friedrich-Alexander-Universität Erlangen-Nürnberg, September 2015.
- [Slac] Slackware. “Howtos : Emulators : binfmt\_misc”. [https://docs.slackware.com/howtos:emulators:binfmt\\_misc](https://docs.slackware.com/howtos:emulators:binfmt_misc) [Accessed 20 August 2020].
- [Tane 14] A. S. Tanenbaum and A. Todd. *Rechnerarchitektur: Von der digitalen Logik zum Parallelrechner*. Pearson Studium ein Imprint von Pearson Deutschland, 2014.
- [Tong 14] X. Tong, T. Koju, and M. Kawahito, Eds. *Optimizing Memory Translation Emulation in Full System Emulators*, IBM Tokyo Research Laboratory, IBM Japan, Ltd., February 2014.
- [Tsch 09] A. Tschoeke, Ed. *Linux on IBM System z10*, IBM Deutschland Research & Development GmbH, TU Kaiserslautern, 2009. [http://www.lgis.informatik.uni-kl.de/cms/fileadmin/users/kschmidt/mainframe/documentation2009/Linux\\_on\\_IBM\\_System\\_z.pdf](http://www.lgis.informatik.uni-kl.de/cms/fileadmin/users/kschmidt/mainframe/documentation2009/Linux_on_IBM_System_z.pdf) [Accessed 21 August 2020].
- [Wang 10] Z. Wang, Ed. *COREMU: A Scalable and Portable Parallel Full-system Emulator*, Fudan University, Parallel Processing Institute, August 2010.



- [Whit 20] B. White, S. C. Mariselli, D. B. de Sousa, E. S. Franco, and P. Paniagua. *Virtualization Cookbook for IBM Z Volume 5 - KVM*. IBM Redbooks, 2020.
- [Yang 19] R. Yang. “Docker Multiarch Builds Quick and Easy”. November 2019. <https://renlord.com/posts/2019-11-09-docker-multiarch/> [Accessed 21 August 2020].



# Glossary

**application layer** The layer in the software stack which is responsible for running applications. 1, iii, 1

**bash** Unix shell and command language integrated in all Linux operating systems. 1

**binary** Executable application built in a single file. 1

**CI/CD** Continuous Integration/Continuous Delivery is a method for automated tests in the software development. 1, iii, 1

**configuration** Saved settings for an application or computer program. 1

**container** A standard unit with the minimum of mandatory software. 1

**container engine** The foundation for managing and organizing containers. 1

**containerized** Splitting programs into different services and deploying it connected in many containers as a single application. 1, iii

**CPU** The Central Processing Unit is a main component of computers for executing instructions. 1

**cross-compilation** Method for compiling code for a different computer system or architecture. 1

**deployment** Fast automated setup of a system with applications and configurations. 1

**Docker daemon** Manages and organizes objects inside of the container engine Docker. 1

**Docker image** A built system listed under "docker images" which is runnable. 1

**Dockerfile** Base file with all installations and configurations for a container image. 1

**emulator** Software for emulating different architectures of other systems for running software on it. 1, iii

**Github** Software version control system by Microsoft for free usage. 1

**Go** Programming language much used for container software. 1

**hard disk** The whole operating system with data is written and saved on this component of the computer system. 1

**high availability cluster** Group of systems which is representing one system with the same software for a minimum amount of down-time. 1

**hypervisor** Software that deploys and run multiple guest virtual machines on real hardware. 1

**issue** A bug report because of a failure in the software. 1

**Java** Programming language much used for Enterprise software development. 1

**JSON** The JavaScript Object Notation is an open standard file format for saving data structured. 1

**JVM** Java Virtual Machine enables running Java applications compiled to Java bytecode. 1

**Kubernetes stack** Kubernetes splitted into different layers from the container environment until the application layer. 1

**library** A suite of reusable code inside of a programming language for software development. 1

**Linux** Free operating system developed by different communities and available as different distributions. 1

**Linux kernel** Software by the Linux community containing all important drivers for running the operating system. 1

**mainframe** Large computer which is able to execute millions of transactions in parallel. 1

**memory** The place inside of a computer for saving data before writing on the hard disk. 1

**mounted** Integrated into a system with the mount command. 1

**multi-architecture images** Images based on system builds usable for multiple system architectures. 1

**package** Software offered comprimized and with dependencies to other software by different Linux distributions. 1

- pod** An association of multiple containers with services for one application. 1
- registry** Built container images are registered and maintained there for general usage. 1
- repository** The location in the internet where software packages are downloadable for installations provided by Linux distributions and other providers. 1
- root permissions** Administrator access permissions with all privileges for a computer system based on Linux/Unix. 1
- s390x** Mainframe system architecture. 1, iii, 1
- scaling** Distributing processes to different cores of a system and executing there. 1
- shell** Terminal of a Linux/Unix system for entering commands. 1
- Ubuntu** Linux distribution maintained by Canonical. 1
- virtual machine** A virtual machine is a system running in a hypervisor as a separate system on another system. 1
- x86** Architecture of a default PC. 1
- z/OS** Operating system by IBM for Z systems. 1
- Z systems** Mainframe hardware by IBM. 1, iii