



TECHNISCHE HOCHSCHULE NÜRNBERG
GEORG SIMON OHM

Fakultät Informatik

Enablement of Kubernetes Based Open-Source Projects on IBM Z

Bachelorarbeit im Studiengang Informatik

vorgelegt von

Sarah Julia Kriesch

Matrikelnummer 303 6764

Erstgutachter:	Prof. Dr. Ralf-Ulrich Kern
Zweitgutachter:	Prof. Dr. Tobias Bocklet
Betreuer:	M.Sc. Alice Frosi
Unternehmen:	IBM Deutschland R & D GmbH

© 2020

Dieses Werk einschließlich seiner Teile ist **urheberrechtlich geschützt**. Jede Verwertung außerhalb der engen Grenzen des Urheberrechtsgesetzes ist ohne Zustimmung des Autors unzulässig und strafbar. Das gilt insbesondere für Vervielfältigungen, Übersetzungen, Mikroverfilmungen sowie die Einspeicherung und Verarbeitung in elektronischen Systemen.

Angaben des bzw. der Studierenden:

Fakultät: Studiengang:

Semester:

Titel der Abschlussarbeit:

Ort, Datum, Unterschrift Studierende/Studierender

Die Entscheidung über die vollständige oder auszugsweise Veröffentlichung der Abschlussarbeit liegt grundsätzlich erst einmal allein in der Zuständigkeit der/des studentischen Verfasserin/Verfassers. Nach dem Urheberrechtsgesetz (UrhG) erwirbt die Verfasserin/der Verfasser einer Abschlussarbeit mit Anfertigung ihrer/seiner Arbeit das alleinige Urheberrecht und grundsätzlich auch die hieraus resultierenden Nutzungsrechte wie z.B. Erstveröffentlichung (§ 12 UrhG), Verbreitung (§ 17 UrhG), Vervielfältigung (§ 16 UrhG), Online-Nutzung usw., also alle Rechte, die die nicht-kommerzielle oder kommerzielle Verwertung betreffen.

Hiermit ☐ genehmige ich, wenn und soweit keine entgegenstehenden Vereinbarungen mit Dritten getroffen worden sind,

☐ genehmige ich nicht,

von Jahren (0 - 5 Jahren ab Datum der Abgabe der Arbeit),

Ort, Datum, Unterschrift Studierende/Studierender

Kurzdarstellung

Open-Source-Projekte entwickeln Software, die frei verfügbar ist. Diese Communitys lassen automatisierte Tests gegen jede Code-Änderung laufen, um die bester Software-Qualität zu garantieren. Diese Tests sollen auf unterschiedlichen Architekturen laufen können. Es ist schwierig, Software für grundlegende Hardware ohne Zugriff darauf - wegen des hohen Preises, zu testen. Deshalb muss die Architektur S390x für Z Systeme für ausgewählte Open-Source-Projekte emuliert (nachgeahmt) und in die entsprechende CI/CD-Pipeline eingefügt werden. Das sollte mit modernen und schnellen Deployment-Methoden realisiert werden, die in den Emulator QEMU integriert werden. Kubernetes wird als Beispiel-Projekt für containerisierte Projekte mit der Sicht, es als Grundlage für alle darauf aufbauenden Projekte verwendet. Ein weiteres Open-Source-Projekt ist Apache Cassandra, um Tests auf der Anwendungsschicht im Kubernetes-Stack zu repräsentieren.

Die minimalen Systemanforderungen müssen für die Einrichtung innerhalb der CI/CD-Infrastruktur beider Projekte mit einem Blick auf die Minimierung an Festplattenplatz, Memory und CPU-Verbrauch auch analysiert werden. Als letzter Schritt wird die automatische Emulation beider Projekte in die Test-Infrastruktur integriert, so dass diese Projekte für die Architektur von Z Systemen aktiviert ist. Allgemein soll diese Methode für weitere Open Source Projekte in der Zukunft wiederverwendet werden können.

Abstract

Open-source-projects are developing software which is free available. These communities are running automated tests against every code change in order to guarantee the best software quality. These tests should be able to run on different architectures. It is difficult to test software for essential hardware without access because of the high price. Therefore, the architecture s390x for Z systems has to be emulated for chosen open-source-projects and included into their CI/CD pipeline. That should be realized with modern and fast deployment methods integrated into the emulator QEMU. Kubernetes is used as an example project for containerized projects in the point of view as the foundation for all establishing projects. Another open-source project is Apache Cassandra to represent tests on the application layer in the Kubernetes stack.

Minimal system requirements have to be analyzed for the setup inside of the CI/CD infrastructure of both projects with an eye on the minimization of space, memory and cpu usage for deployments, too. As the last step, the automated emulation of both projects will be integrated into the test infrastructure, that these projects are enabled for the architecture of Z systems. Overall, this method should be reapplied for further open-source-projects in the future, too.

Contents

1. Introduction	1
1.1. Mainframe Computers	2
1.2. Hardware Emulation	2
1.3. Open Source Projects	2
1.3.1. Kubernetes	2
1.3.2. Apache Cassandra	3
2. Emulation	5
2.1. QEMU	5
2.2. System Emulation	6
2.3. User Mode Emulation	6
2.4. Emulation of Different Architectures	6
2.4.1. Binfmt_misc	7
2.4.2. Qemu-user-static	7
2.4.3. BuildX	8
2.5. File Systems	8
2.5.1. UnionFS	9
2.5.2. Ext4	10
2.5.3. Mounting an External File System	11
3. Prerequisites	13
3.1. Registration of Qemu-S390x	13
3.2. Enabling Multi-Architecture Images	14
3.3. Cross-Compilation of a S390x Linux Kernel on x86	14
3.4. Optimized QEMU Command	15
4. Continuous Integration	17
5. Kubernetes	19
5.1. Overview	19
5.2. Multi Staging Dockerfile	19
5.3. Installation	20
5.4. Installation of the Ultimate Go	21

5.5. Building Tests	22
5.5.1. E2e.test	23
5.6. Run in the Main Dockerfile	23
5.7. Integration into CI/CD	23
6. Cassandra	25
6.1. Overview	25
6.2. Deployment	25
6.3. Workaround Because of an JVM Issue During the Build	26
6.4. Icinga Monitoring Check for Tests	28
6.5. Integration into CI/CD	28
7. Outlook	29
8. Summary	31
A. Supplemental Information	33
List of Figures	35
List of Tables	37
List of Listings	39
Bibliography	41
Glossary	43

Chapter 1.

Introduction

One business introduced by IBM is the mainframe, now known as a Z system. It is possible to run Linux on it. There is a large community behind Linux and open-source projects. Open source does not contain only Linux. There are different applications and other software developed by open-source communities. The mainframe hardware architecture is different from the architecture of a home PC. The Z system architecture is called s390x and a default system x86. Z systems are costly, and not every open-source community can pay such a system. Most open-source Contributors have got systems with the architecture x86 at home. Therefore, it should be possible to test hardware dependencies for s390x on x86.

Our focus is on Kubernetes-based open-source projects in this Bachelor Thesis. Some Kubernetes-based projects should be emulated for Z systems in the CI/CD test infrastructure by these open-source projects. That will be done on systems by these communities. As the first step, the emulator will be chosen with the focus on functionality for Z systems on x86 architecture. After that, Kubernetes is installed in a Docker container. Tests should be able to be run on this system, too. That will be integrated into the emulation environment for an automated start. The CI/CD system should be able to execute all tests then.

The same will be done with the NoSQL database Cassandra for the Apache community to represent the whole system stack from Kubernetes until the application layer for container platforms. Other points are minimal systems requirements and minimal systems sizes. Here are different methods evaluated to minimize the system for emulation. The goal of this Bachelor Thesis is to offer emulated Z systems for different open-source projects to test their software for hardware dependencies, so that it allows to release new versions in the CI/CD pipeline of the respective project for running on the hardware architecture s390x without the available hardware. Deployments of the latest software version on Github and running tests have to be automated for that.

1.1. Mainframe Computers

Mainframe computers are large computers. Some of them are part of the Z series¹ by IBM. They are not only used as internet servers. They can handle large numbers of transactions in one second[Tane 14, p.56]. Such Z systems do not use the well known x86 architecture. They are built with s390x. IBM has developed this architecture. The current architecture has been introduced in late 2000 and is supported by the Linux kernel since late 1999 [Bloc 19, p.15]. IBM Z contains IBM Z pervasive encryption for comprehensive protection around the data on the system[Lasc 20, p.4]. Such systems are offering a horizontal and upright scaling of processes, which allows the operation of many hundred virtual systems in parallel[Tsch 09, p.13]. The traditional operating system for mainframes has been z/OS. Z systems have been optimized for open-source software as Linux[Lasc 20, p.8]. The goal by IBM has been to offer a combination of a robust and securable hardware platform with the power of different Linux distributions. Linux is used as a base operating system for this Bachelor Thesis.

1.2. Hardware Emulation

Not everybody has access to expensive hardware or hardware with essential architecture. The software should be able to run on most relevant hardware architectures. The solution for Software Developers is hardware emulation. You can test based on hypervisors with the hardware emulation whether the software is running correctly. So you can run different operating systems and applications for specialized hardware in virtualization software. It is possible to enable other hardware architectures than the host has got. That will be done with the implementation of a VM on a host system with a different target architecture[Rose 15, p.3]. In our case, it should be possible to run applications for mainframes with the target architecture s390x on a host system with the architecture x86.

1.3. Open Source Projects

1.3.1. Kubernetes

Kubernetes² is an open-source project for container orchestration, well known as k8s. This project was started by Google. A Kubernetes cluster has at least one Master node and one Worker node for high availability. The Master node is responsible for managing all Worker nodes with applications. All configurations and distributions are performed from there. New Worker nodes are added to the Kubernetes cluster with "join" on the Master node, too.

¹<https://www.ibm.com/it-infrastructure/z/hardware/>

²<https://kubernetes.io/>

The Worker node is receiving containers with different services. Every Worker node can communicate with other Worker nodes in the cluster.

Kubernetes is scalable for using containers as distributed services in a pod. A pod is representing something as a single server splitted into different connected containers with all services for a running application. Pods can be replicated to multiple nodes for high availability. Kubernetes is configurable with different container runtimes, as Docker³, containerd⁴ or CRI-O⁵ for example. The Container Runtime Interface (CRI) is necessary for managing container images, the life cycle of container pods, networking and help functions [Scho 19, p.16]. The mainly used container runtime is Docker in the Kubernetes project. The content of a container is described in the format of a Dockerfile. This format is supported by most container runtimes.

1.3.2. Apache Cassandra

Apache Cassandra⁶ is a NoSQL database management system by the Apache Foundation. The project had been started internally at Facebook and has been released as an open-source-project in 2008. Cassandra provides continuous availability, high performance, and linear scalability besides offering operational simplicity and effortless replication across data centers and geographies[Data]. It is recommended for mission-critical data. The Cassandra Query Language (CQL) is equal to SQL and includes JSON support. CQL includes an abstraction layer that hides implementation details of the structure.

NoSQL databases store are using other methods than relational database management systems for data access and have got a different structure. Cassandra is using hashing to fetch rows from the table as a column-oriented database management system (DBMS). Such a DBMS stores data tables by column rather than by row.

NoSQL databases were developed as massively scalable database management systems that are able to write and read data anywhere while distributing availability to billions of users. This field is well known as Big Data.

³<https://www.docker.com/>

⁴<https://containerd.io/>

⁵<https://cri-o.io/>

⁶<https://cassandra.apache.org/>

Chapter 2.

Emulation

2.1. QEMU

QEMU is an open-source emulator available in most Linux distributions. It is embedded in different virtualization tools as KVM and XEN, too. So QEMU is well tested and has got all the necessary features for emulations. Additionally, you can emulate other architectures on different hardware. QEMU is a generic emulator and can emulate different system architectures. It can also be used for emulation of obsolete hardware[Opsa 13, p.24]. It has been extended for various architectures as x86, ARM, PowerPC, Sparc32, Sparc64, MIPS and s390x. That is a processor emulator using binary translation[Butt 11] which is executing and translating emulated code based on blocks. Each block has got one entry and one exit point[Wang 10, p.5]. The binary translation will be executed with the "Tiny Code Generator" (TCG) inside of QEMU. That is a small compiler replacing GCC because of unlimited releases and changes with it. The TCG is converting the blocks of target instructions into a default intermediate form. Subsequently, it has to be compiled for the host or target architecture. If a binary for a new target architecture is necessary, the frontend of TCG will be converted while QEMU is ported to a new architecture. The TCG integrates new code for the new host architecture in the background then. It is also dedicated to improve performance with avoiding repeated translations by buffering already translated code[Cota 17]. The TCG takes care of the emulation of the guest processor running as a thread launched by QEMU. The memory of a guest is allocated during launch and is mapped into the address space of the QEMU process[Opsa 13, p.29].

It is possible to run unmodified guest operating systems. The open-source projects can use any Linux distribution as their base operating system then because QEMU is integrated as a package as default. QEMU does not emulate the whole hardware. That is only possible for the CPU. Therefore, QEMU is used for emulations in this Bachelor Thesis.

2.2. System Emulation

The System Emulation emulates a whole system with hardware, the operating system (with the kernel) and the userspace (with application processes). The system (hardware and operating system) will be translated unmodified. The benefit of system emulation (compared to user mode emulation) is that you can run privileged instructions[Butt 11]. Additionally, it can be used as an application development platform where specific hardware is not available. System emulation with emulated hardware is slower than a real machine because instructions executed directly in the guest hardware have to be emulated in software. That implies multiple host instructions because of the translation for a single guest instruction as a result[Tong 14, p.1]. It is possible to reduce the supported and attached additional devices with special options. The deactivation of a graphic card can be specified with **-nographic** as an example. Afterwards, less hardware has to be emulated, which has given the result of better performance.

System emulation benefits from the virtualization support as KVM. So most CPU operations are not required to emulate.

2.3. User Mode Emulation

The User Mode Emulation does not emulate the whole system. It is possible to reproduce application processes in QEMU with a minimal system for a specific application. This emulation type is working on a syscall level. Therefore, the application has to be runnable as a process executable itself. An external Linux kernel will be built, and the application can be mounted via a loaded Docker image in a hard disk image. The emulator is using the Linux kernel to emulate system signal calls then. That can be managed by mapping system calls of the target to an equivalent system call on the host. The user-mode emulation can run directly non-privileged instructions or is using system calls to ask for a selected service from the operating system[Butt 11].

2.4. Emulation of Different Architectures

It is achievable to emulate different architectures on another hardware architecture. The package qemu-user-static has to be installed then, and the preferred architecture has to be registered in binfmt. binfmt_misc is a kernel module. You can register other architectures within that, so multiple other architectures can be run on this host. Not only QEMU can be used for system emulation. Container technologies as Docker include an integrated QEMU compatibility as an additional emulation feature for building images for other architectures.

This technology has got the name BuildX. Accordingly, a hybrid virtualization approach is practicable with different virtualization technologies as with QEMU and Docker together.

2.4.1. Binfmt_misc

Binfmt_misc is a feature in the kernel which allows invoking almost every program by simply typing its name in the shell. That permits to execute user space applications such as emulators and virtual machines for other hardware architectures. It recognises the binary-type by matching some bytes at the beginning of the file with a magic byte sequence. These executable formats have to be registered in the file `/proc/sys/fs/binfmt_misc/register`. The structure of the configuration for the registration is the following:

:name:type:offset:magic:mask:interpreter:flags

The **name** is the name of the architecture for the binary format. The **type** can be **E** or **M**. E is for executable file formats as `.exe` for example. M is used for a format identified by a **magic** number at an absolute **offset** in the file and the **mask** is a bitmask indicating which bits in the number are significant[Slac] (which is used in our case). The offset can be kept empty. The magic is a byte sequence of hexadecimal numbers. The **interpreter** is the program that should be invoked with the binary[Gün]. The path has to be specified for it. The field **flags** is optional. It checks different aspects of the invocation of the interpreter. The **F** flag is necessary in our case for "fix binary". That supposes that a new binary has to be spawned if the misc file format has been invoked.

On that way, it is possible to register another system architecture (s390x on x86) on the host system and it is realizable to emulate the other architecture afterwards.

The necessary command for s390x registration can be found under 3.1. The package **binfmt-support** has to be installed first for using this kernel feature.

2.4.2. Qemu-user-static

The package **qemu-user-static**¹ enables the execution of multi-architecture container emulation based on **binfmt_misc** with **QEMU**. This package by various Linux distributions includes a set of binfmt configurations for various architectures together with an amount of statically compiled QEMU emulators that other architectures can run on another architecture[Yang 19]. When building for another architecture, static binaries are important because of the possibility of an "exec init error" without them.

The installation contains emulators for all available architectures supported by the QEMU project incl. s390x aligned with the specific host architecture. In this manner described, you can build and emulate for different architectures on the same host (see 3.4 Optimized Qemu Command). That is the base for BuildX by Docker, too.

¹<https://github.com/multiarch/qemu-user-static>

2.4.3. BuildX

As mentioned above, Docker is using QEMU for emulations with BuildX². That is an integrated "experimental Feature" since version Docker 19.03. That implies that it is not enabled as default because it is very new and for experiments. It has to be activated with `DOCKER_CLI_EXPERIMENTAL=enabled` (see 3.1 Registration of Qemu-S390x). The disadvantage of the provided BuildX inside of Docker is that you do not receive the up-to-date version of BuildX. There is a version with a more detailed output and better working with an additional **buildx** behind **docker build**. This version can be cloned from <https://github.com/docker/buildx> and installed with **make install**.

In general, BuildX is a Docker CLI plugin for expanding the "docker build" command. That includes the possibility of multi-architecture builds and exceptional output configuration. That involves not only the upload possibility to your local docker registry (see the output of **docker images**). That includes the output of "docker build" (which creates a docker image based on the requirements and installations in the used Dockerfile) into a local directory, a tar archive or a public registry on Docker Hub. Docker Hub is the most used and maintained registry for Docker images.

The **Dockerfile** is using the requirements by the base docker image in the FROM command on top of the Dockerfile together with installations and configurations executed by all RUN commands or attached scripts with the command ADD. The **-t** flag is adding a special name as a tag under **docker images**. The special architecture for multi-architecture builds can be specified with the additional option **--platform**. The option **build create** provides the option to create different build instances for multiple build combinations of architectures.

2.5. File Systems

Docker is using the Union File System which is not compatible with QEMU. QEMU can integrate only hard disks formatted for default Linux file systems as ext2, ext3, ext4, XFS and Btrfs. The driver virtio_blk is used to mount external file systems and emulates read/write in the physical block device[Barb18]. Following it is possible to integrate and start nonnative systems in QEMU. Docker is advantageous to setup and start systems fast. So it would be nice to integrate the docker file system into QEMU. After building a docker image, it is realizable to export the file system into a directory with the name **rootfs** with the command `docker export $(docker create initrd-s390x) | tar -C "rootfs"-xvf -`.

Linux has got the feature that it is adventitious to reformat directories for file systems and to copy/ mount content into this one. Reformatting the default docker file system UnionFS

²<https://github.com/docker/buildx>

to another one as ext4 for example can be done then.

QEMU is accepting the new file system as a block device for the guest system then. The default path to the mounted file system as a hard disk is `/dev/vda` as the first partition[[Whit 20](#), p.22].

2.5.1. UnionFS

Docker does not use any default Linux file system. Docker images are based on the Union File System (UnionFS)[[Ashr 15](#), p.21]. This file system is using different file system layers with grouping directories and files in branches. The first layer is the typical Linux boot file system with the name `bootfs`. That performs the same as in a Linux virtualization stack with using memory at first and unmounting to receiving RAM free by the `initrd` disk image. So the `bootfs` can be used inside of another Linux file system to mount in an virtualization stack for a successful boot process with QEMU. The next layer contains the base image with the operating system given by the "FROM" command. Then every docker command inside of the Dockerfile is adding an additional layer with the installation of applications or building binary files. That is the reason that every executed docker command is receiving an own id in the disk memory during the build process. Every separate docker command is using his own disk space. So a docker image can grow really fast. It is reasonable to compress so much as possible of different routines into one docker command. All sizes of different docker layers will exist continuously inside of the new file system.

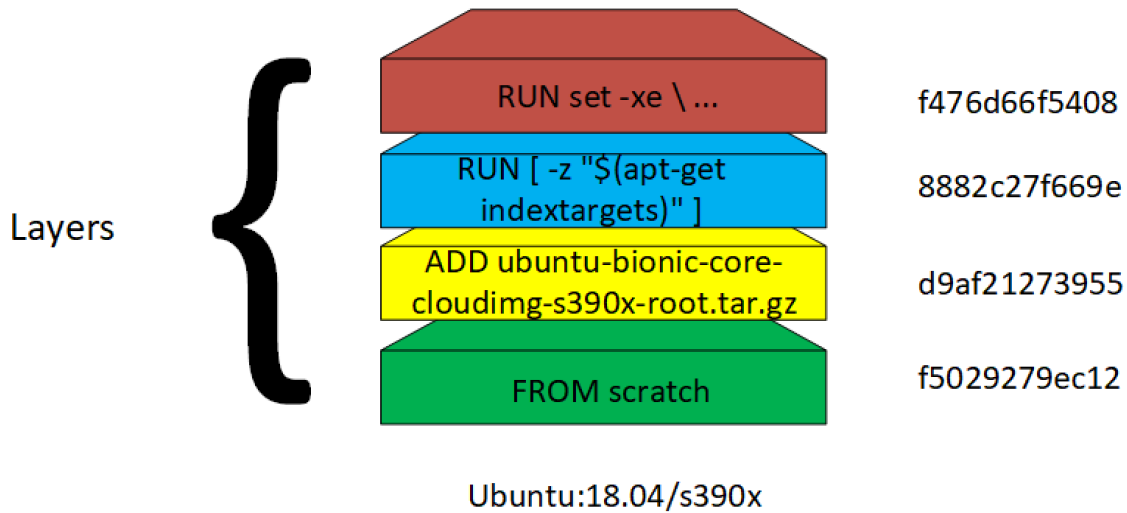


Figure 2.1.: Union File System

2.5.2. Ext4

Ext4 is a journaling file system (the same as ext3). The journal is registering transactional all changes on the operating system with meta data. So no data are lost after a system failure. They can be restored based on the journal without a save procedure by the user. Such journal file systems in the ext file system family are working with blocks[Seuf 15, p.20]. The journal is splitted into the journal super block, descriptor blocks, commit blocks and revoke blocks. The super block contains all meta data of the journal. Descriptor blocks include the special destination adress and a sequence number. Commit blocks are available to flag logged transactions. Revoke blocks are flagging blocks not logged any more.

All meta data from the journal can be relocated in inodes because changes of file system meta data are registered, too.

The difference to the journal in ext3 are the option writing asynchronous commit blocks and additional check sums for journal transactions[Seuf 15, p.28].

Ext4 provides a better performance than ext3. This file system can be formatted with `mkfs.ext4` which is available in every Linux distribution. This tool is creating a group descriptor table for further group descriptors what allows an expanding of the file system. The file system can grow a maximum of the multiple 1024 of his existing size because of the

saved space[Seuf15, p.21]. Flex groups are used in ext4 for merging different block groups to one logic block group. All data from inodes are only saved in the first flex group then. So the search for files is more powerful because meta data are allocated at one place.

Another feature of ext4 is the possibility of inline files and inline directories. So small files and directories can be saved directly in inodes instead of data blocks. From this follows less disk space consumption[Seuf15, p.24].

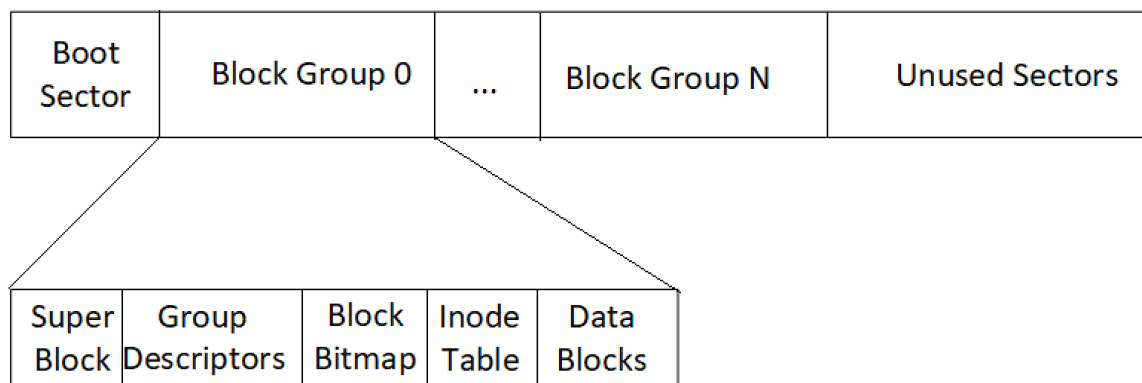


Figure 2.2.: Ext4

2.5.3. Mounting an External File System

In 2.5 "File Systems" is described how to export a Docker file system into a directory. That has got the format of the 2.1 UnionFS at this moment. QEMU contains the beneficial tool **qemu-img** for creating loadable images of file systems. These require the minimal size of the docker image. The output of the size is given by `docker images | grep ${id}`. QEMU understands only rounded numbers. Therefore, that has to be rounded up to a full number. Cassandra is using the command `qemu-img create -f raw cassandra.img ↪ 2G` (see 6.2 Cassandra Deployment) for a Docker image size about 1.2G then. Firstly, `cassandra.img` is empty. The content of **rootfs** has to be transferred into this image file. But QEMU does not know UnionFS as a file system. That is the reason for formatting `cassandra.img` with `mkfs.ext4 -F cassandra.img` to ext4. Writing into `cassandra.img` is only executable if the image is mounted into a directory under `/mnt/` with `mount -o loop ↪ cassandra.img /mnt/rootfs`. In the next step the content of the file system directory can be copied into `/mnt/rootfs/` with `cp -r rootfs/* /mnt/rootfs/..`. Then all data of the Docker image are saved into `cassandra.img` because it is mounted into `/mnt/rootfs/`. The image `cassandra.img` is usable for deployments with **qemu-system-s390x** (see 6.2 Run Cassandra) at the end.

Chapter 3.

Prerequisites

Different software is necessary to run qemu or docker for multiple architectures. Therefore, docker and qemu should be installed. Additionally, `qemu-user-static`¹ and `binfmt_misc`² are important for running multi-architecture containers.

Root permissions are required for installation, configuration and running emulated systems.

It is possible to install packages as `binfmt-support` and `qemu-user-static` by different Linux distributions, but it is recommended to use the ultimate version for s390x.

3.1. Registration of Qemu-S390x

The packages `qemu-user-static` and `binfmt-support` should be installed because of the configuration.

The Linux kernel module `binfmt_misc` can be mounted with the following command:

```
mount binfmt_misc -t binfmt_misc /proc/sys/fs/binfmt_misc
```

Ultimate stable releases of `qemu-user-static` can be found under <https://github.com/multiarch/qemu-user-static/releases/>. The release v5.0.0-2 is used for the project and downloaded with `wget https://github.com/multiarch/qemu-user-static/releases/download/v5.0.0-2/x86_64_qemu-s390x-static.tar.gz` for the specific version of `qemu-s390x-static` on x86. That is extracted to the directory `/usr/bin/` with the command `sudo tar -xvzf x86_64_qemu-s390x-static.tar.gz -C /usr/bin/` then. The architecture name before `qemu` is the architecture of the host system. The name between `qemu` and `static` is the emulated architecture. This archive contains only the new version for the unique architecture. The configuration will be used by the installed `qemu-user-static`.

¹<https://github.com/multiarch/qemu-user-static>

²<https://www.kernel.org/doc/html/latest/admin-guide/binfmt-misc.html>

s390x binaries have to be registered for s390x. That is done with the following commands:
`sudo -i` and

```
echo ':qemu-s390x:M::
\x7fELF\x02\x02\x01\x00\x00\x00\x00\x00\x00\x00\x00\x02\x00\x16:
\xff\xff\xff\xff\xff\xff\xff\x00\xff\xff\xff\xff\xff\xff\xff\
\xff\xfe\xff\xff:/usr/bin/qemu-s390x-static:OCF' > /proc/sys/fs/binfmt_misc/
register
```


Listing 3.1: Register s390x binaries

3.2. Enabling Multi-Architecture Images

Docker is configured to build only for the own architecture, which is x86 at open-source projects. The "Experimental" flag exists for new available features which are not ready for production. So you can build docker images for s390x on x86 then. That can be added with the following:

```
{
"experimental": enabled
} >> /etc/docker/daemon.json
```

Listing 3.2: Docker Experimental Flag

The configuration of the Docker daemon is written in the JSON format. That is saved in `/etc/docker/daemon.json`. After a restart of the docker daemon, there should be listed this experimental flag in the command `docker version`. An alternative way is to export this flag as an environment variable in the shell with `export DOCKER_CLI_EXPERIMENTAL=enabled` to enable it. Now it is possible to build docker images for s390x on x86 with `docker build`  `--platform=linux/s390x -t image-example:s390x .` based on a Dockerfile in the existing directory.

3.3. Cross-Compilation of a S390x Linux Kernel on x86

QEMU needs a built Linux kernel to start a system. These packages are necessary:

bison, flex, gcc, gcc-s390x-linux-gnu, libssl-dev

One kernel source code for a certain version has to be downloaded from kernel.org with the following command then: `wget https://git.kernel.org/torvalds/t/linux-5.7-rc7.tar.gz`. It can be extracted with `tar -xf linux-5.7-rc7.tar.gz`. The Makefile is in the directory `linux-5.7-rc7`. Therefore the command `make ARCH=s390 defconfig localyesconfig` has

to be executed to create the default configuration for s390x and `make ARCH=s390 CROSS_COMPILE=s390x-linux-gnu- -j6` for the compilation. The kernel has the name `bzImage` in the directory `arch/s390/boot/` then and can be used in the `qemu` command.

3.4. Optimized QEMU Command

Every additional device requires additional performance and time for starting the system. So the systems requirements had to be figured out to be minimal for every given open-source project and for running tests on it. That counts for the number of CPUs, too.

The kernel option is receiving the path to the built s390x kernel with the name `bzImage`. In the given 3.3 "Optimized QEMU command" the `qemu` command is executed in the directory with the built Linux kernel and the path is unimportant.

The option **-m** is available to add the minimal guest memory matching the system requirements of every open-source project. **-nodefaults** is deactivating default additional devices activated in QEMU. Only the console is necessary for receiving an output and debugging. So that is added as a device. Cassandra as a project does not need any network interface or parallelism. The option **-nographic** is responsible for not adding any graphical interface. So we save system requirements. The option **-smp** is the minimal number of CPUs for the guest. The file system of containers can be loaded as a hard disk with the option **-hda** which is explained in every chapter of the Bachelor Thesis for a considered open-source project. That is the ideal option to mount a minimal file system for every application or system. `/dev/vda` is the partition name and `rdinit` is used for using `bash` as a default shell.

This command can be used for user mode emulation with Apache Cassandra as an example.

```
/
usr/bin/qemu-system-s390x -kernel bzImage -m 4G -M s390-ccw-virtio
    ↪ -nodefaults -device sclpconsole,chardev=console -parallel none -net
    ↪ none -chardev stdio,id=console,signal=off,mux=on -mon
    ↪ chardev=console -nographic -smp 3 -hda /data/cassandra.img --append
    ↪ 'root=/dev/vda rw console=ttyS0 rdinit=/bin/bash'
```

Listing 3.3: Optimized QEMU Command

Chapter 4.

Continuous Integration

Hier kommt ein Teil zu CI/CD

Chapter 5.

Kubernetes

5.1. Overview

Kubernetes is a container platform for high availability clusters. There exist plugins for integration tests for Kubernetes with the name kubetest¹. They contain conformance tests, as well as e2e tests (end-to-end), too. That can be all built and executed on the system. Therefore, a Dockerfile for setting up Kubernetes and building tests with Go is necessary. The challenge is, that 2 big Github repositories have to be cloned and integrated into the docker image. That is using a lot of space. The solution is using a multi staging Dockerfile. So 2 different Dockerfiles are used in one Dockerfile and one of them is used for building. The other one is used for the installation and testing with built tests. At the end the size of the docker image has got only the size of the test image regardless of the repository size in the mother Dockerfile.

5.2. Multi Staging Dockerfile

A multi-staging Dockerfile is using different systems in one Dockerfile for different stages. These systems are receiving special names as indicators with "AS" behind the "FROM" with the base image name. Default this feature is used for building applications in one stage and executing the copied application in another stage. The same counts for cloning Github repositories and building binary files based on it. On this way, a lot of space is saved. Concluding, the docker image has got only the size of the executing system with the application file (without all the code). That is an "experimental feature" at the moment. Therefore the **experimental flag** is necessary to export or set before using it (see 3.2 Multi-Architecture Images).

¹<https://kubetest.readthedocs.io/en/latest/>

5.3. Installation

Kubernetes needs a lot of packages for running and for tests. That will be all installed with the RUN command. apt.kubernetes.io has got later versions of Kubernetes than the Ubuntu repository. Therefore this repository has to be added to Ubuntu. kub-build is the name of the mother Dockerfile to be able to copy needed files and directories from there.

In the intallation part of the Dockerfile the Kubernetes repository <https://apt.kubernetes.io> for Debian packages has to be registered together with with the gpg key used by <https://packages.cloud.google.com/apt/doc/apt-key.gpg>. A Dockerfile is installing only necessary packages. Therefore, apt-transport-https, apt-utils, curl, git, ca-certificates, gnupg-agent and software-properties-common have to be installed first for the following installation. The system requires the latest update with `apt-get update` after the registration of the additional repository besides of the default Ubuntu repositories imported with the base Ubuntu image "s390x/ubuntu:18.04" in the FROM command. After that the packages docker.io, kubelet and kubeadm can be installed from kubernetes.io. **Docker.io** contains the container engine docker with all docker commands. CRI/O or containerd would be allowed, too. The Docker daemon will be used because that is the main used container engine of the Kubernetes project and all tests are running with it.

Kubelet is the primary node agent running on each node. He is responsible that different containers can run together in a pod. Pods are deployable units defined in JSON or a yaml file. They include one or a group of containers with shared storage and network resources. Hosting of distributed systems with different services in different containers can work together. Consequential one pod is something as one "logical host".

Kubeadm is the administration tool to setup clusters. It is necessary to upgrade Kubernetes to other versions, too. Clusters can be initialized. The network can be configured and the command **kubectrl** (Kubernetes Control Plane) for adding additional nodes to a cluster can be initialized.

`apt-mark hold` is keeping these special versions of kubelet, kubeadm and kubectrl.

```

FROM s390x/ubuntu:18.04 AS kub-build

# The author
MAINTAINER Sarah Julia Kriesch <sarah.kriesch@ibm.com>

#Installation
RUN echo "Installing necessary packages" && \
apt-get update && apt-get install -y \
apt-transport-https \
apt-utils \
systemd \
curl \
git \
ca-certificates \
gnupg-agent \
software-properties-common \
&& curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add - \
&& echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" \
> /etc/apt/sources.list.d/kubernetes.list \
&& apt-get update && apt-get install -y \
docker.io \
kubenet \
kubeadm \
&& apt-mark hold kubenet kubeadm kubectl \
&& apt-get clean \
&& rm -rf /var/lib/apt/lists/* /tmp/* /var/tmp/* \
&& systemctl enable docker

```

Figure 5.1.: Kubernetes Installation

5.4. Installation of the Ultimate Go

There were some issues with older Go versions as 1.10 during building tests for Kubernetes. Therefore a higher version (min. 1.13) should be used. It is recommended to use the ultimate go version for last versioned Kubernetes tests. It is possible to receive the version number of the ultimate go release with the command `curl https://golang.org/VERSION?m=text`. This version number has to be included before `linux-s390x.tar.gz` for downloading the special s390x archive from the go directory by dl.google.com. Then the version number has to be called with curl inside of another curl command with the whole path to the special tar

archive on dl.google.com. Every tar archive has got the same structure for every version (`$version.$architecture.tar.gz`). On this way the latest version of Go is integrable into the curl command that it can be installed. Directories for bin, pkg and src have to be created after extracting this tar archive in the `/root/` directory. They are not integrated in the tar archive.

The environment variables for GOROOT, GOPATH and PATH have to be set with ENV on the top of the Dockerfile for successful builds later. PWD is added because Github repositories have to be cloned to this directory.

The ENV variables will be on the top of the Dockerfile. The part for the "Installation of Go" will be attached to the end of the Kubernetes installation part.

```
ENV GOROOT=/root/go
ENV GOPATH=/root/go
ENV PATH=$GOPATH/bin:$PATH
ENV PATH=$PATH:$GOROOT/bin
ENV PWD=/root/go/src/

#Installation of latest GO
&& echo "Installation of latest GO" && \
curl "https://dl.google.com/go/ \
$(curl https://golang.org/VERSION?m=text).linux-s390x.tar.gz" \
| tar -C /root/ -xz \
&& mkdir -p /root/go/{bin,pkg,src}
```

Figure 5.2.: Go Installation

5.5. Building Tests

After a successful installation of go, it is possible to build and install the Kubernetes test environment. At first the directory `k8s.io` has to be created because kubernetes-tests are looking for this directory as a mother directory. The repository `test-infra` by the Kubernetes project has to be cloned to there. The repository **test-infra** contains all tests for Kubernetes provided by the community. These can be used of course and are updated continuously. That is the reason to clone this repository inside of the Dockerfile. Inside of this `test-infra` directory `kubetest` can be installed with **go install**. That is downloading all available Kubernetes-Tests. So you can use them to test the own Kubernetes cluster and the used software.

The name of the most relevant tests for the Kubernetes community is "conformance tests". These conformance tests are executed with `e2e.test` which can be built with `make` inside of the `kubernetes` repository. Therefore this repository has to be cloned to `k8s.io`, too. These tests certify the software to comply regular standards. Only with complying these standards, Kubernetes software is allowed to become Kubernetes certified².

```
&& cd $PWD \  
#Clone test-infra  
&& mkdir -p $GOPATH/src/k8s.io \  
&& cd $GOPATH/src/k8s.io \  
&& git clone https://github.com/kubernetes/test-infra.git \  
/root/go/src/k8s.io/test-infra \  
&& cd /root/go/src/k8s.io/test-infra/ \  
#Install kubetest  
&& GO111MODULE=on go install ./kubetest \  
#Build test binary  
&& git clone https://github.com/kubernetes/kubernetes.git \  
/root/go/src/k8s.io/kubernetes \  
&& cd /root/go/src/k8s.io/kubernetes/  
  
CMD make WHAT="test/e2e/e2e.test vendor/github.com/onsi/ginkgo/ginkgo cmd/kubect1"
```

Figure 5.3.: Test Building for Kubernetes

5.5.1. E2e.test

5.6. Run in the Main Dockerfile

5.7. Integration into CI/CD

²<https://github.com/cncf/k8s-conformance>

Chapter 6.

Cassandra

6.1. Overview

6.2. Deployment

IBM is providing a Dockerfile¹ for Apache Cassandra especially for Z systems with the ultimate version of Cassandra on Github. This file can be cloned to the system and the Docker image will be built with the command `docker build buildx --platform=linux/s390x --squash -t cassandra:s390x .` in the directory with the Cassandra Dockerfile for the architecture s390x.

Squash is an option to compress a Docker image and combine commands in a Dockerfile automatically. The prerequisites for building s390x images on x86 are set during the emulation preparation. The command `docker images` has to show the registered Docker image with the name `cassandra:s390x` then. It should be possible to integrate this Docker image into the `qemu` command. Therefore, a `qemu-image` will be created with an rounded given size besides of the Docker image in the `docker images` command. So the command `qemu-img create -f raw ↪ cassandra.img 2G` can be used. This image needs any Linux file system because QEMU does not know the Docker file system. The image is formatted with the command `mkfs.ext4 ↪ -F cassandra.img` then. For receiving the file system of the docker image a directory with the name `rootfs` has to be created and the command `docker export $(docker create ↪ cassandra:s390x) | tar -C "rootfs"-xvf -` is exporting the docker image into the directory `rootfs`. Following transfers the content of `rootfs` into the image `cassandra.img`.

```
mkdir /mnt/rootfs
mount -o loop cassandra.img /mnt/rootfs
cp -r rootfs/* /mnt/rootfs/.
```

Listing 6.1: Mount rootfs

¹<https://github.com/linux-on-ibm-z/dockerfile-examples/tree/master/ApacheCassandra>

Now it is possible to run the system with Cassandra:

```
/
usr/bin/qemu-system-s390x -kernel bzImage -m 40G -M s390-ccw-virtio
    ↪ -nodefaults -device sclpconsole,chardev=console -parallel none -net
    ↪ none -chardev stdio,id=console,signal=off,mux=on -mon
    ↪ chardev=console -nographic -smp 3 -hda /data/dockerfile-examples/
    ApacheCassandra/cassandra.img --append 'root=/dev/vda rw console=ttyS0
    ↪ rdinit=/bin/bash'
```

Listing 6.2: Run Cassandra

6.3. Workaround Because of an JVM Issue During the Build

Apache Cassandra is using OpenJDK 8 for running as default. There are some discussions² about supporting OpenJDK 11 at the moment. The Dockerfile mentioned above is using AdoptOpenJDK 8 at the build process is working without any problems on Z systems. There exists the following JVM issue during the build process on x86 which is the reason for non successful builds of Java applications there:

²<https://lists.apache.org/thread.html/r38f6beaa22e247cb38212f476f5f79efdc6587f83af0397406c06d7c%40dev.cassandra.apache.org>


```

#
# A fatal error has been detected by the Java Runtime Environment:
#
# SIGILL (0x4) at pc=0x0000004012539240, pid=18, tid=21
#
# JRE version: OpenJDK Runtime Environment (11.0.7+10) (build 11.0.7+10)
# Java VM: OpenJDK 64-Bit Server VM (11.0.7+10, mixed mode, tiered, compressed oops,
g1 gc, linux-s390x)
# Problematic frame:
# J 1 c1 java.lang.Object.<init>()V java.base@11.0.7 (1 bytes) @ 0x0000004012539240
[0x0000004012539200+0x0000000000000040]
#
# Core dump will be written. Default location: //core
#
# An error report file with more information is saved as:
# //hs_err_pid18.log
Compiled method (c1)      943      1      3      java.lang.Object::<init> (1 bytes)
  total in heap [0x0000004012539010,0x00000040125393b8] = 936
  relocation    [0x0000004012539178,0x00000040125391a8] = 48
  constants     [0x00000040125391c0,0x0000004012539200] = 64
  main code     [0x0000004012539200,0x0000004012539300] = 256
  stub code     [0x0000004012539300,0x0000004012539358] = 88
  metadata      [0x0000004012539358,0x0000004012539370] = 24
  scopes data   [0x0000004012539370,0x0000004012539380] = 16
  scopes pcs    [0x0000004012539380,0x00000040125393b0] = 48
  dependencies  [0x00000040125393b0,0x00000040125393b8] = 8
Compiled method (c1)      948      1      3      java.lang.Object::<init> (1 bytes)
  total in heap [0x0000004012539010,0x00000040125393b8] = 936
  relocation    [0x0000004012539178,0x00000040125391a8] = 48
  constants     [0x00000040125391c0,0x0000004012539200] = 64
  main code     [0x0000004012539200,0x0000004012539300] = 256
  stub code     [0x0000004012539300,0x0000004012539358] = 88
  metadata      [0x0000004012539358,0x0000004012539370] = 24
  scopes data   [0x0000004012539370,0x0000004012539380] = 16
  scopes pcs    [0x0000004012539380,0x00000040125393b0] = 48
  dependencies  [0x00000040125393b0,0x00000040125393b8] = 8
Could not load hsdis-s390x.so; library not loadable; PrintAssembly is disabled
#
# If you would like to submit a bug report, please visit:
#   https://github.com/AdoptOpenJDK/openjdk-support/issues
#
Aborted (core dumped)

```

Figure 6.1.: JVM Build Issue

This case has been tested with AdoptOpenJDK 8 and AdoptOpenJDK 11. The library `hsdis-s390.so` is not loadable for both Java versions on x86, but during the build process on s390x. This issue does not exist with the version AdoptOpenJDK 14, but Apache Cassandra does not work with this version.

That is an emulation bug of qemu which has to be fixed. This issue has been reported to KVM Developers inside of IBM.

There is a workaround for the development process for finishing the project. The Docker image is buildable on s390x. Therefore, the Docker image has been built for s390x on s390x with `docker build --squash -t cassandra ..`. Afterwards, this image can be saved as a tar archive with `docker save cassandra > cassandra.tar` on the host. It has to be transferred with `rsync` to the x86 system then. The image inside of the tar archive can be included to all existing images of `docker images` with the command `docker load ↪ --input cassandra.tar`. After this action, the Docker image is usable as it would have been built on this system.

6.4. Icinga Monitoring Check for Tests

Icinga is an open-source monitoring system. There are many checks for different services available. One monitoring check exists for Apache Cassandra. It is testing whether Java would be working, the service Apache Cassandra would be up and running and whether CQL commands would be possible to execute without any issues.

This monitoring check is written in Perl and is a nice choice for automated tests with error messages in the case if any output would not give an **OK**.

6.5. Integration into CI/CD

Chapter 7.

Outlook

Ich kommt ein Teil zur Zukunft zum Projekt

Chapter 8.

Summary

Hier kommt eine Zusammenfassung

Appendix A.

Supplemental Information

Zusatzinformation

List of Figures

2.1. Union File System	10
2.2. Ext4	11
5.1. Kubernetes Installation	21
5.2. Go Installation	22
5.3. Test Building for Kubernetes	23
6.1. JVM Build Issue	27

List of Tables

List of Listings

3.1. Register s390x binaries	14
3.2. Docker Experimental Flag	14
3.3. Optimized QEMU Command	15
6.1. Mount rootfs	25
6.2. Run Cassandra	26

Bibliography

- [Ashr 15] W. Ashraf. *The Docker EcoSystem*. Gitbook, 2015.
- [Barb 18] D. H. Barboza. “Enhancing QEMU virtio-scsi with Block Limits vital product data (VPD) emulation”. October 2018. <https://developer.ibm.com/technologies/linux/articles/enhancing-qemu-virtio-scsi-with-block-limits-vpd-emulation/#sec3.1> [Accessed 17 August 2020].
- [Bloc 19] B. Block, Ed. *Modern Mainframes & Linux Running on Them*, IBM Deutschland Research & Development GmbH, Chemnitzer Linux-Tage, March 2019. <https://chemnitzer.linux-tage.de/2019/media/programm/folien/173.pdf> [Accessed 07 February 2020].
- [Butt 11] K. Butt, A. Qadeer, and A. Waheed. “MIPS64 User Mode Emulation: A Case Study in Open Source Software Engineering”. In: *2011 7th International Conference on Emerging Technologies*, pp. 1–6, 2011.
- [Cota 17] E. Cota, P. Bonzini, A. Bennée, and L. Carloni. “Cross-ISA machine emulation for multicores”. In: *2017 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, pp. 210–220, February 2017.
- [Data] Datastax. “What is Cassandra?”. <https://www.datastax.com/cassandra> [Accessed 26 August 2020].
- [Gün] R. Günther. “Kernel Support for miscellaneous Binary Formats (bintfmt_misc)”. <https://www.kernel.org/doc/Documentation/admin-guide/bintfmt-misc.rst> [Accessed 22 August 2020].
- [Lasc 20] O. Lascu, B. White, J. Troy, and F. Packheiser. *IBM z15 Technical Instruction*. IBM Redbooks, 2020.
- [Opsa 13] J. M. Opsahl. *Open-source virtualization - Functionality and performance of Qemu/KVM, Xen, Libvirt and VirtualBox*. Master’s thesis, University of Oslo, May 2013.
- [Rose 15] D. S. Rosenthal. *Emulation & Virtualization as Preservation Strategies*. Andrew W. Mellon Foundation, 2015.

- [Scho 19] B. Scholl, T. Swanson, and P. Jausovec. *Cloud Native: Using Containers, Functions, and Data to Build Next-Generation Applications*. "O'Reilly Media, Inc.", 2019.
- [Seuf 15] S. Seufert. *Implementierung eines forensischen Ext4-Incode-Carvers*. Master's thesis, Friedrich-Alexander-Universität Erlangen-Nürnberg, September 2015.
- [Slac] Slackware. "Howtos : Emulators : binfmt_misc". https://docs.slackware.com/howtos:emulators:binfmt_misc [Accessed 20 August 2020].
- [Tane 14] A. S. Tanenbaum and A. Todd. *Rechnerarchitektur: Von der digitalen Logik zum Parallelrechner*. Pearson Studium ein Imprint von Pearson Deutschland, 2014.
- [Tong 14] X. Tong, T. Koju, and M. Kawahito, Eds. *Optimizing Memory Translation Emulation in Full System Emulators*, IBM Tokyo Research Laboratory, IBM Japan, Ltd., February 2014.
- [Tsch 09] A. Tschoeke, Ed. *Linux on IBM System z10*, IBM Deutschland Research & Development GmbH, TU Kaiserslautern, 2009. http://wwlgis.informatik.uni-kl.de/cms/fileadmin/users/kschmidt/mainframe/documentation2009/Linux_on_IBM_System_z.pdf [Accessed 21 August 2020].
- [Wang 10] Z. Wang, Ed. *COREMU: A Scalable and Portable Parallel Full-system Emulator*, Fudan University, Parallel Processing Institute, August 2010.
- [Whit 20] B. White, S. C. Mariselli, D. B. de Sousa, E. S. Franco, and P. Paniagua. *Virtualization Cookbook for IBM Z Volume 5 - KVM*. IBM Redbooks, 2020.
- [Yang 19] R. Yang. "Docker Multiarch Builds Quick and Easy". November 2019. <https://renlord.com/posts/2019-11-09-docker-multiarch/> [Accessed 21 August 2020].

Glossary

application layer The layer in the software stack which is responsible for running applications. 1, iii, 1

bash Unix shell and command language integrated in all Linux operating systems. 1

binary Executable application built in a single file. 1

CI/CD Continuous Integration/Continuous Delivery is a method for automated tests in the software development. 1, iii, 1

configuration Saved settings for an application or computer program. 1

container A standard unit with the minimum of mandatory software. 1

container engine The foundation for managing and organizing containers. 1

containerized Splitting programs into different services and deploying it connected in many containers as a single application. 1, iii

CPU The Central Processing Unit is a main component of computers for executing instructions. 1

cross-compilation Method for compiling code for a different computer system or architecture. 1

deployment Fast automated setup of a system with applications and configurations. 1

Docker daemon Manages and organizes objects inside of the container engine Docker. 1

Docker image A built system listed under "docker images" which is runnable. 1

Dockerfile Base file with all installations and configurations for a container image. 1

emulator Software for emulating different architectures of other systems for running software on it. 1, iii

Github Software version control system by Microsoft for free usage. 1

Go Programming language much used for container software. 1

hard disk The whole operating system with data is written and saved on this component of the computer system. 1

high availability cluster Group of systems which is representing one system with the same software for a minimum amount of down-time. 1

hypervisor Software that deploys and run multiple guest virtual machines on real hardware. 1

issue A bug report because of a failure in the software. 1

Java Programming language much used for Enterprise software development. 1

JSON The JavaScript Object Notation is an open standard file format for saving data structured. 1

JVM Java Virtual Machine enables running Java applications compiled to Java bytecode. 1

Kubernetes stack Kubernetes splitted into different layers from the container environment until the application layer. 1

library A suite of reusable code inside of a programming language for software development. 1

Linux Free operating system developed by different communities and available as different distributions. 1

Linux kernel Software by the Linux community containing all important drivers for running the operating system. 1

mainframe Large computer which is able to execute millions of transactions in parallel. 1

memory The place inside of a computer for saving data before writing on the hard disk. 1

mounted Integrated into a system with the mount command. 1

multi-architecture images Images based on system builds usable for multiple system architectures. 1

package Software offered comprimized and with dependencies to other software by different Linux distributions. 1

- pod** An association of multiple containers with services for one application. 1
- registry** Built container images are registered and maintained there for general usage. 1
- repository** The location in the internet where software packages are downloadable for installations provided by Linux distributions and other providers. 1
- root permissions** Administrator access permissions with all privileges for a computer system based on Linux/Unix. 1
- s390x** Mainframe system architecture. 1, iii, 1
- scaling** Distributing processes to different cores of a system and executing there. 1
- shell** Terminal of a Linux/Unix system for entering commands. 1
- Ubuntu** Linux distribution maintained by Canonical. 1
- VM** A Virtual Machine (VM) is a system running in a hypervisor as a separate system on another system. 1, 2
- x86** Architecture of a default PC. 1
- z/OS** Operating system by IBM for Z systems. 1
- Z systems** Mainframe hardware by IBM. 1, iii, 1